

Unwanted Traffic: Denial of Service Attacks

Dan Boneh

What is the D S?

◆ G a : a e a a g e e i h i e c ing

◆ H : **Amplification**

- S a n b e f a c e = big effec

◆ T e f a i f i c a i n a c :

- D S b g :

◆ De igh t a a ing ne achine di a
e ice

- D S f d :

◆ C and b -ne gene a e f d f e e

DDoS can happen anywhere

◆ This occurs:

- Same DDoS attack type (botnet):

- ◆ LIP
- ◆ TCP/UDP
- ◆ Application

- DDoS mitigation

◆ Sad fact:

- Content providers are not designed to handle DDoS attacks

Waveform of 802.11b DS-SS

◆ Radiating a packet : initialization of channel.

◆ Packet of DS-SS [Bard, Savage, 03]

▪ NAV (Network Allocation Vector):

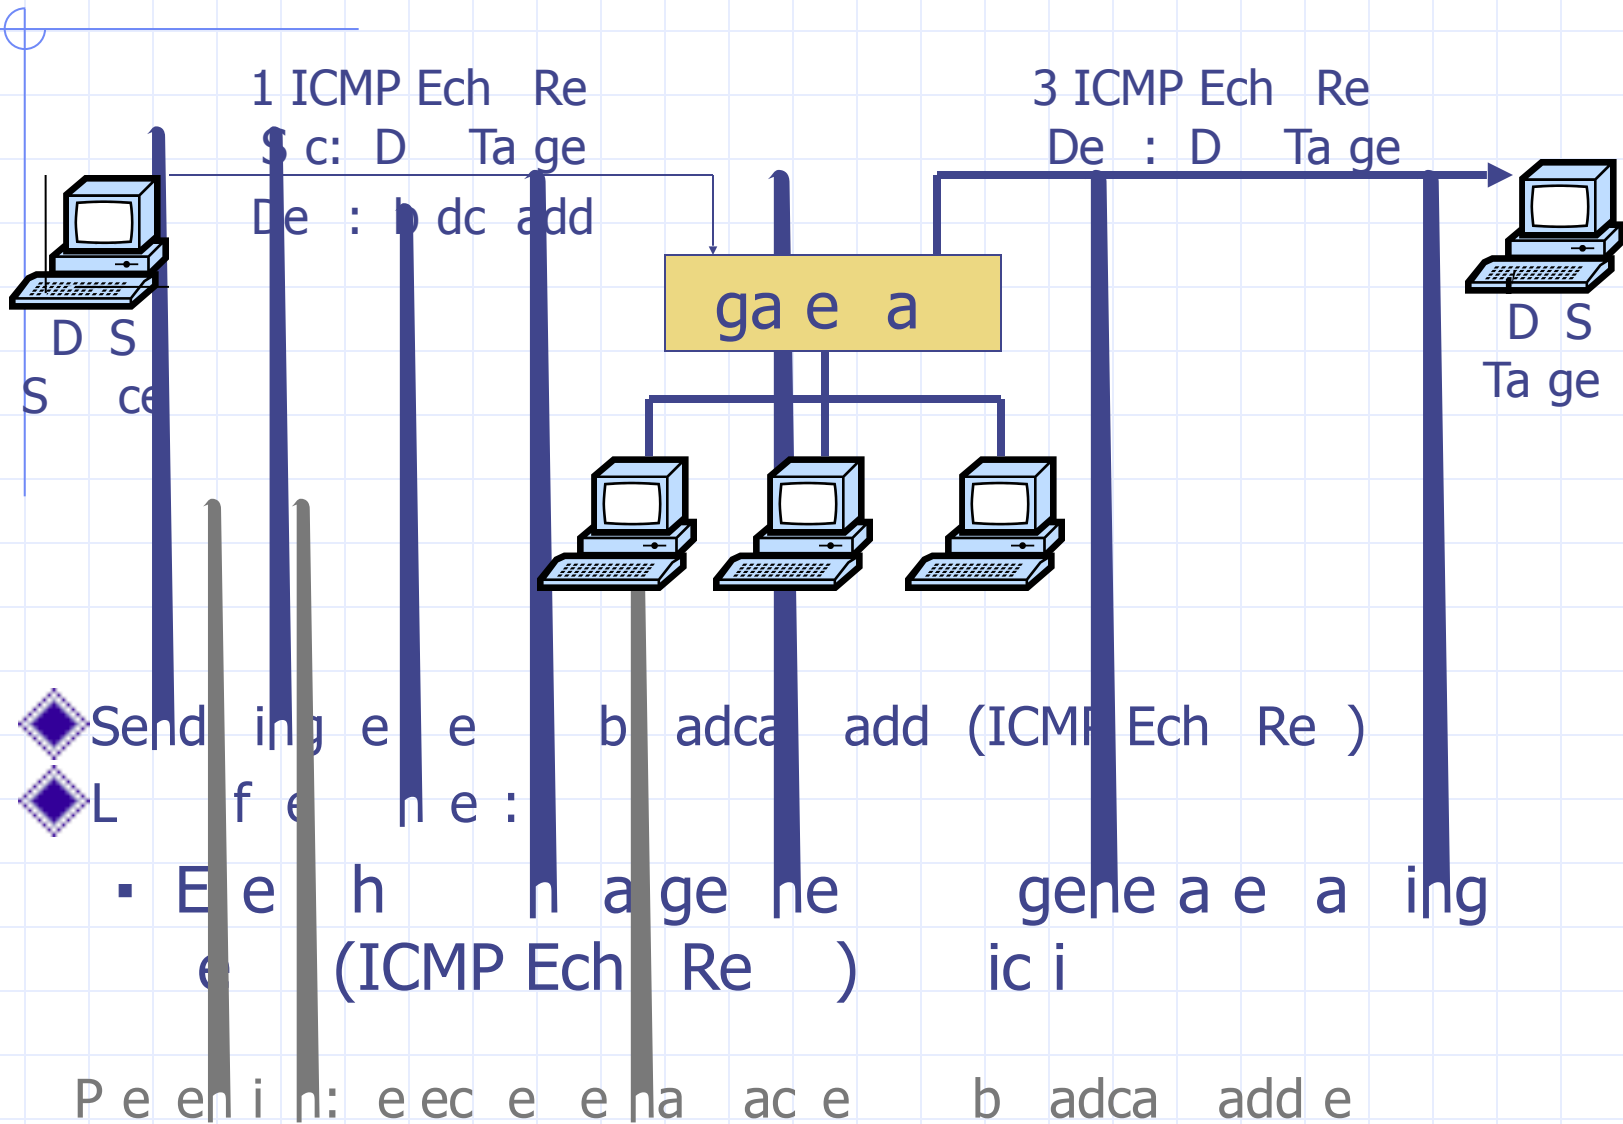
- ◆ 16-bit field. Max value is 2767
- ◆ All nodes can receive channel for NAV encoding
- ◆ Nodes can send and receive NAV encoding
- ◆ NAV is a 16-bit field for 802.11b channel

▪ De-activation in b/g:

- ◆ All nodes can send data to AP
- ◆ Deactivation in b/g
- ⇒ a node can be deactivated and then




Significati in DNS



M de n da e a e (Ma 13

DNS A ifica i n a ac : (×50 a ifica i n)



2006: 0.58M en e e n n e n e (Ka in -Shiff a)

2017: 15M en e e (en e e ec . g)

⇒ 3/2013: DD S a ac gene a ing 309 Gb f 28 in .

Scale, Targeting and Frequency of Attacks





Figure 13
Source: Arbor Networks, Inc.

Peak DDoS Attack Size (January 2010 to Present)



Source: Arbor Networks, Inc.

Feb. 2014: 400 Gbps via NTP amplification (4500 NTP servers)

Review : IP Header format



Connection

- Unreliable
- Best effort

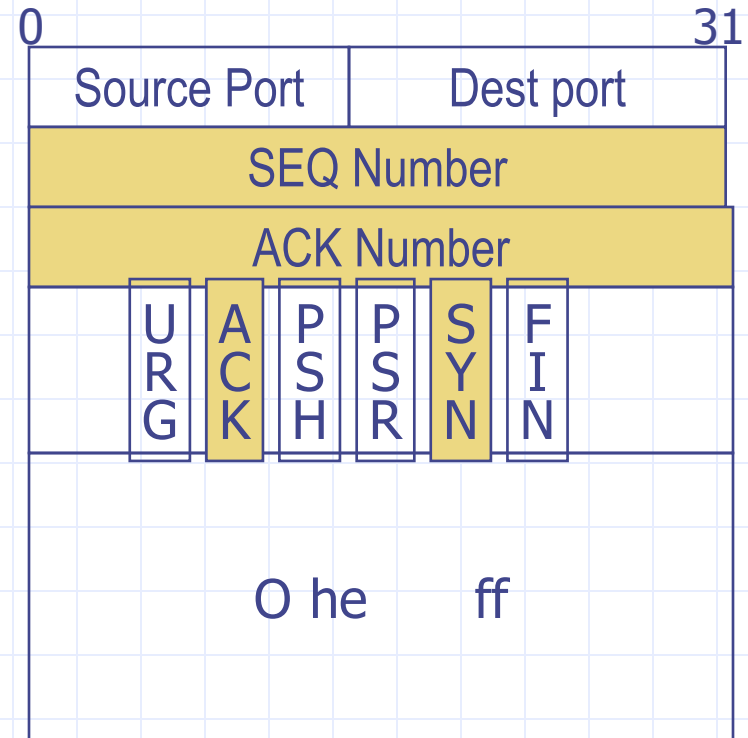
0	31
Version	Header Length
Type of Service	
Total Length	
Identification	
Flags	Fragment Offset
Time to Live	
Protocol	
Header Checksum	
Source Address of Originating Host	
Destination Address of Target Host	
Options	
Padding	
IP Data	

Review: TCP Header




TCP:


- Sequence based
- Congestion control
- In-order delivery



Review : TCP Handshake



TCP SYN Flood I:



Single machine:

SYN Packets with
**random source IP
addresses**

For each packet, the server
sends back a SYN-ACK packet.
Since the source IP addresses are
random, the server never receives
the connection request.

SYN Flood (March 48, 1996)

OS	Backlog size
Linux 1.2.x	10
FreeBSD 2.1.5	128
WinNT 4.0	6

Backlog size : 3 min

- Attacker needs only 128 SYN packets every 3 minutes
- Low rate SYN flood

La SYN et la définition

The β :
the efficiency ()
before and

Nh- in:
▪ Increase in efficiency
decrease in

Cec in (head):
▪ **Syncookies**:
▪ Self-ance head

Secure


[Beinlein, Schen]

- ◆ Idea: encode the data in the sequence number (SN)
- ◆ Server sends Client in SYN-ACK connection:
 - $T = 5\text{-bit counter incremented every } 64 \text{ sec}$
 - $M = \text{MAC}_e(\text{SAddr}, \text{SP}, \text{DAddr}, \text{DP}, \text{SN}_C, T)$ [24 bit]
 - ◆ e : encrypted and digest
 - $\text{SN}_S = (T \cdot L)$ ($L = 24 \text{ bit}$)
 - **Server does not save state** (like TCP in a server)
- ◆ Hence client sends in ACK ($\text{AN} = \text{SN}_S$, $\text{SN} = \text{SN}_C + 1$)
 - Server accepts the connection if $\text{AN} = \text{SN}_S$

SYN flood : background [MVS 01]


◆ SYN initiated by the attacker ⇒ SYN/ACK and then

Bacula e ea e en



Ethiopia access

(ATLAS 07)



◆ Attacks detected:

- 15 ICMP flood, 4 TCP SYN flood

◆ Bandwidth

- 12 attacks: **70-95 Mbps for over 10 hours**

◆ Attacks affecting access inside Ethiopia

- Ethiopia is in:

- ◆ Ethiopian ISP blocked a foreign traffic in
accessed

⇒ DSA had internet access inside Ethiopia

Mailbox (e.g. Mail 9/2016 in Kib)

C and b a f d ecific a ge : (DD S)

- F d i h SYN, ACK, UDP, and GRE ac e
- 623 Go (ea) f $\approx 100K$ c i ed I T de ice
- A eo i e:
 - Sa a e ne in ne e
 - Rand ce IP \Rightarrow

a ac SYN

he a e a ea SYN

- Wha d ???

Country	% of Mirai botnet IPs
Vietnam	12.8%
Brazil	11.8%
United States	10.9%
China	8.8%
Mexico	8.4%
South Korea	6.2%
Taiwan	4.9%
Russia	4.0%
Romania	2.3%
Colombia	1.5%

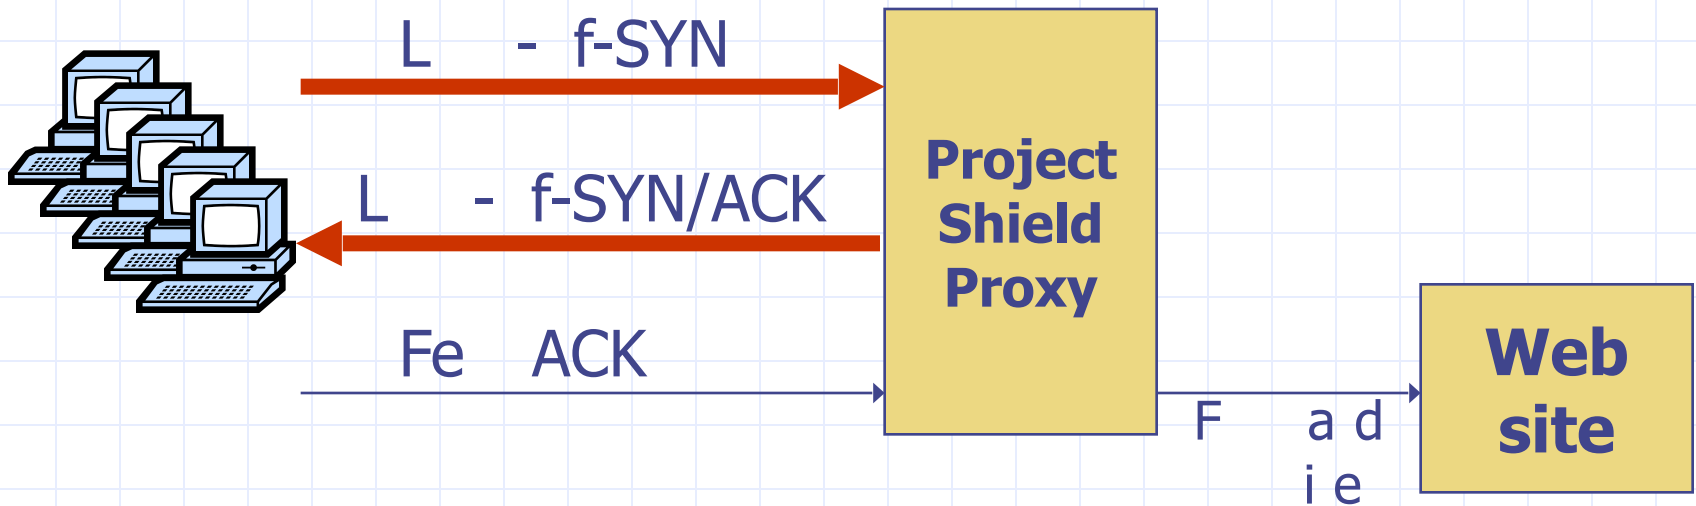
Figure 3: Top countries of origin of Mirai DDoS attacks

Google Shield

Project Shield is a non-profit organization.

(Co-located services: AdSense, Analytics, etc.)

Idea: to provide a distributed TCP connection pool



Single access: GET method

◆ Client and browser:

- Client establishes TCP connection to web site
- Send http GET request
- Receive


◆ Web browser sends request

◆ browser:


- Access can be made and receive IP.
- ◆ Receive can be file
- Page can be displayed.

A ea - d e a e: G H b (3/2015)


Ja a c i -ba ed DD S:




gi h b.c



ne
end e



in ec
i ageF d



a
e e

i ageF d.

```
f nc i i f d()
a TARGET = ' ic i - eb i e.c /inde . h ?
a and = Ma h.f (Ma h. and () * 1000)
a ic = ne I age()
ic c = 'h ://'+TARGET+ and+'= a '
e In e a(i gf d, 10)
```

W d HTTPS
e en hi DD S?

DNS DoS Attack

(e.g. Dyna 10/2016)

DNS runs on UDP port 53

- DNS server forwards requests to authoritative DNS servers

DNS Amplification

- Flood DNS servers with requests
- **Random source IP address** in UDP packets
- Take advantage of DNS servers (create additional traffic)

Dyn attack: Distributed MITM-based

- Attacked 100,000+ servers in

⇒ Dyn cannot handle large volume of DNS requests

⇒ Distributed denial of service (DDoS) attacks, e.g.,

D S Mi iga i n

1. Challenge

◆ Idea: find x on a machine

◆ Model we had before:

- Given challenge C find X such that

$$SHA-1(SB_n(SHA-1(C || X))) = 0^n$$
- Attack on n : a few executed 2^n times
- For $n=64$ a few about .3 sec on 1GHz machine
- Main problem: checking the idea.

◆ Doing Disassembly:

- Estimate the bits in the machine
- When on a machine: don't see the idea

E a e

◆ GET f d (RSA 99)

- E a e change: $C = T \oplus e - e - n$
- Fi da a ac e c n a e i n
 - ◆ O he i e TCP c n n e c i i e d

◆ SSL hand ha e D S: (SD 03)

- Cha e n e C ba e d n TLS e i n ID
- Se e : chec e i n bef e RSA dec .

Benefit and Efficiency

- ◆ **Hardness of change:**
 - Decided based on DSA action.
- ◆ **Efficiency:**
 - Rejected change both efficient and effective
 - However, the efficiency is a secondary indicator
 - ◆ Efficiency alone is not enough and alone cannot connect

◆ Meeting point

- Main e acce e ni :

- Be e e :

- ◆ D -G dbe g-Nb, C 03

- 29

2. CAPTCHA

◆ Idea: we if have a connection if a human



- ◆ A nice application is a DDOS [Kip 05]
- During an attack: generate CAPTCHA and send it to the user if a request is received.
 - Prevent the CAPTCHA to be sent to the IP address.


3. Să se identifice în

Gă: identifice

Uiaegă: b c a ac a he ce

1. Incoming

(RFC 2827, 3704)



I e en a i n be

ALL ISP d n i . Re i e g ba .

- If 10% f ISP d n i e en \Rightarrow n defen e
- N incen e f de en

2017:

- 33% f A . S e a e f fab e
(fe .caida. g)
- 23% f ann nced IP add e ace i fab e

Reca : 309 Gb a ac ed n 3 ne (3/2013)

2. Tacebac

[Sa age e a . 00]

◆ G a :

- Gi en e f a ac ac e
- De e i e a h ce

◆ H : change e ec d inf in ac e

◆ A i n :

- M e e ain nc i e
- A ac e end an ac e
- R e f a ac e ic i e ain e a i e
ab e

Si e e n d



Wie a h i n n e a c e

- Each e add i n IP add e ac e
- Vic i ead a h f ac e



P b e :

- Re i e e a c e i n a c e


◆ Pa h can be n g

◆ N e a f i e d i n c e n IP f a

- Change ac e f a ch e ec

Be the idea

- ◆ DD Simple and an
ac e n a e a h
- ◆ S e n e i n e a c h
ac e
- Each e
pabi i ca e
n add e
- Fi ed a e e g a d e
f a h e n g h



Edge Saving



Da a field in the edge structure:

- Edge start and end IP addresses
- Distance: number of hops since edge reached



Making a record of the R (if a link is in the head of the table) then


if $i \in R$ then add e
 if $i \in 0$ in distance field
 if $i \in R$ then end
 if distance == 0 then end
 increment distance field

Edge Sampling: iterative



Package received


- R_1 receives package from the sender
- Package contains address, end, distance



Edge Sampling: iterative

◆ Begin with edge

- R_1 chooses a new edge
- Set distance 0



Edge Saving



Finishing edge

- R_2 chosen as the next edge
- Distance is 0

◆ When end of edge, increase distance 1

ac e	R_1	R_2	1
------	-------	-------	---



Edge Saving



Increment distance

- R_3 cannot be reached by edge
- Distance > 0

◆ Increment distance 2



Path decomposition

◆ Each vertex appears in at most k bags

◆ Bag B_i contains vertices v_1, \dots, v_{d_i}

- Each (v_i, v_j) induces an edge

◆ # bags needed to cover G is $\chi(G)$

$$E(X) < \frac{\ln(d)}{p}$$

$$p(1-p)^{d-1}$$

he e i a ing babi i , d i eng h f a h

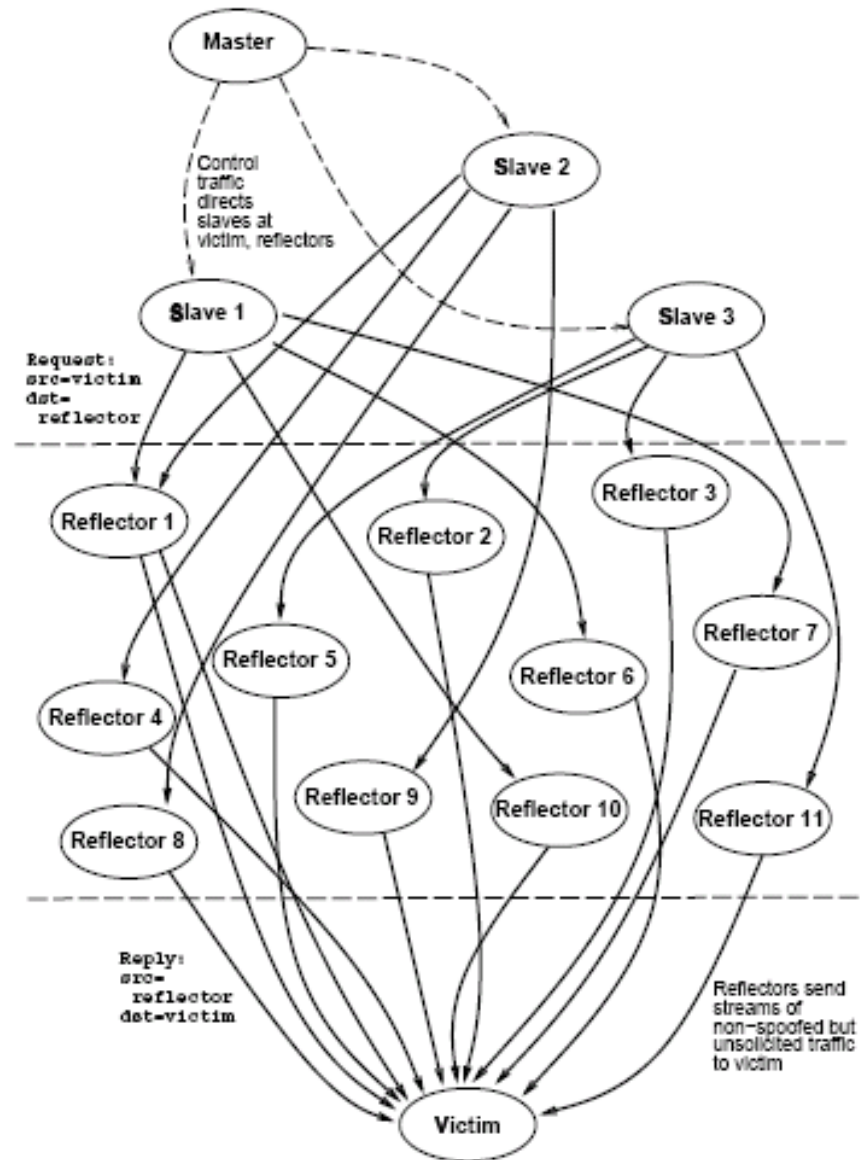
D S A ac

◆ Single Master

◆ Master generates flood

◆ Zi i n f ef ec
hide b

- Ki acebac and
hbac e h d



Ca abi i ba ed defe n e

Capabilities based defense

◆ Ande, R., & Wehe. .

- Peer engineering media - facilities for capabilities . SIGCOMM 04.

◆ Yao, Perrig, and Song.

- Siff: A secure neighbor discovery algorithm for i iga e DD S
ding a ac IEEE S&P 04.

◆ Yang, Wehe, Ande .

- A D S-i ing ne a chi ec e.
SIGCOMM 05

Ca abii ba ed defen e

◆ Ba icidea:


- Re ei e can ecif ha ac e he an

◆ H :

- Sende e e ca abii in SYN ac e
 - ◆ Pa h den ife ed i i # e f ne ce
- Re ei e e nd i h ca abii
- Sende inc de ca abii in a f e ac e
- **Main point:** R e n f a d:
 - ◆ Re e ac e , and
 - ◆ Pac e i h a id ca abii

Canability of a definition

- Canability can be extended if the following
 - Block access to the



Taken Home Message:

- ◆ Definition of Service a social action:
 - Must be considered as designed initiative
- ◆ Satisfaction:
 - In the individual's hands DD Satisfaction
 - Managerial actions in : Client's, Admin, Financial, etc.
- ◆ Managerial actions designed

