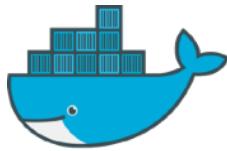


CS155: Infrastructure Security Building Blocks



@diogomonica





- Infrastructure Company
- Running on 10% of the internet
- Infra: Go, some python



- Mobile payments company
- Moving \$60 billion annually
- Infra: Java & Ruby, some Go

Agenda

- 1. Why infrastructure security?**
- 2. Biased timeline of how we got here**
- 3. What is Docker**
- 4. Docker security tetrodinos**
- 5. Demo**

Why Infrastructure Security?

~1980

2000



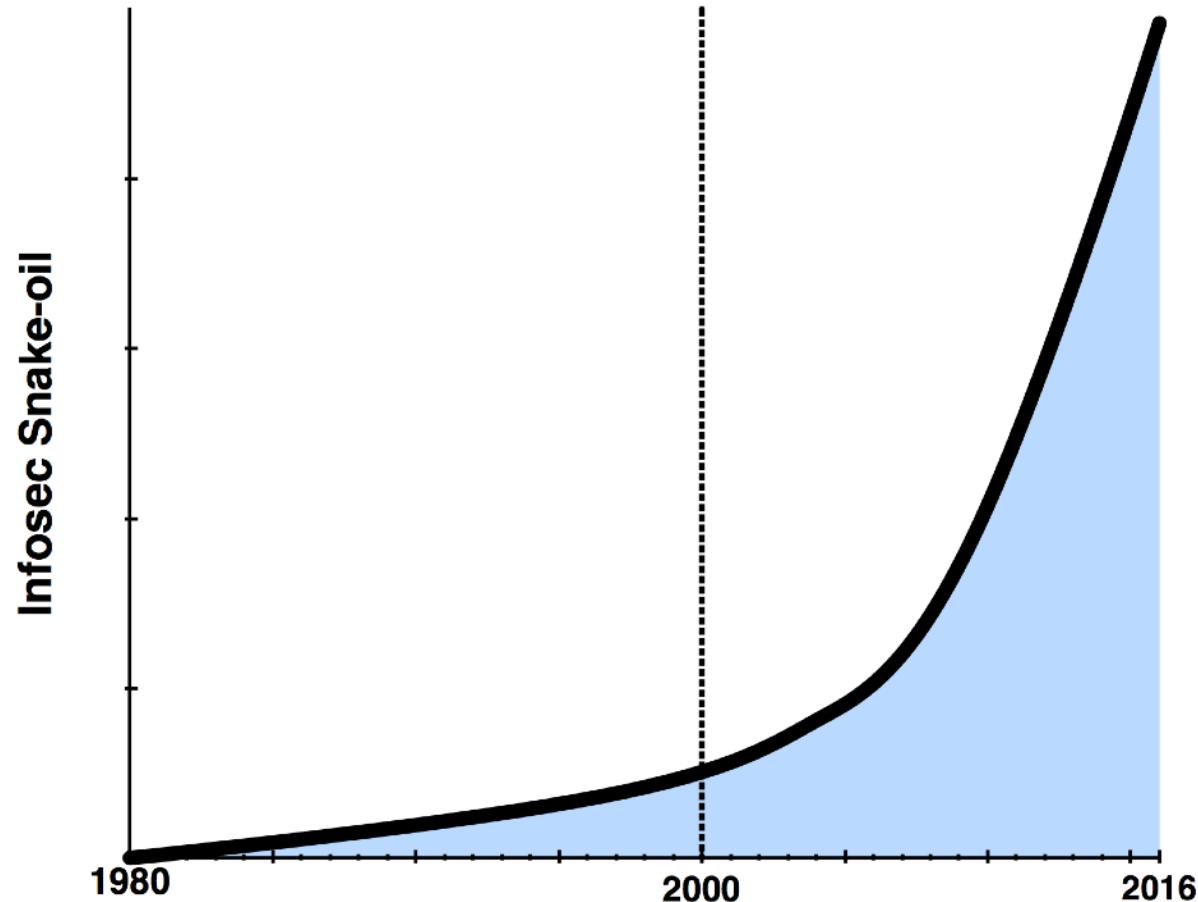
The Golden Ages

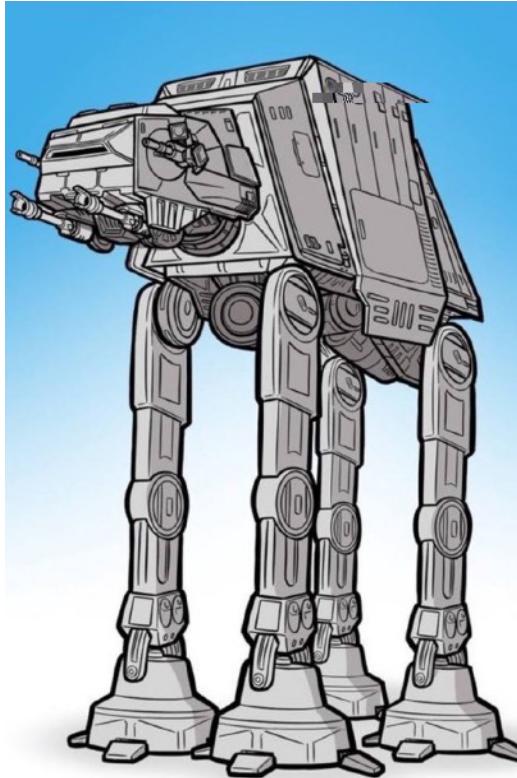


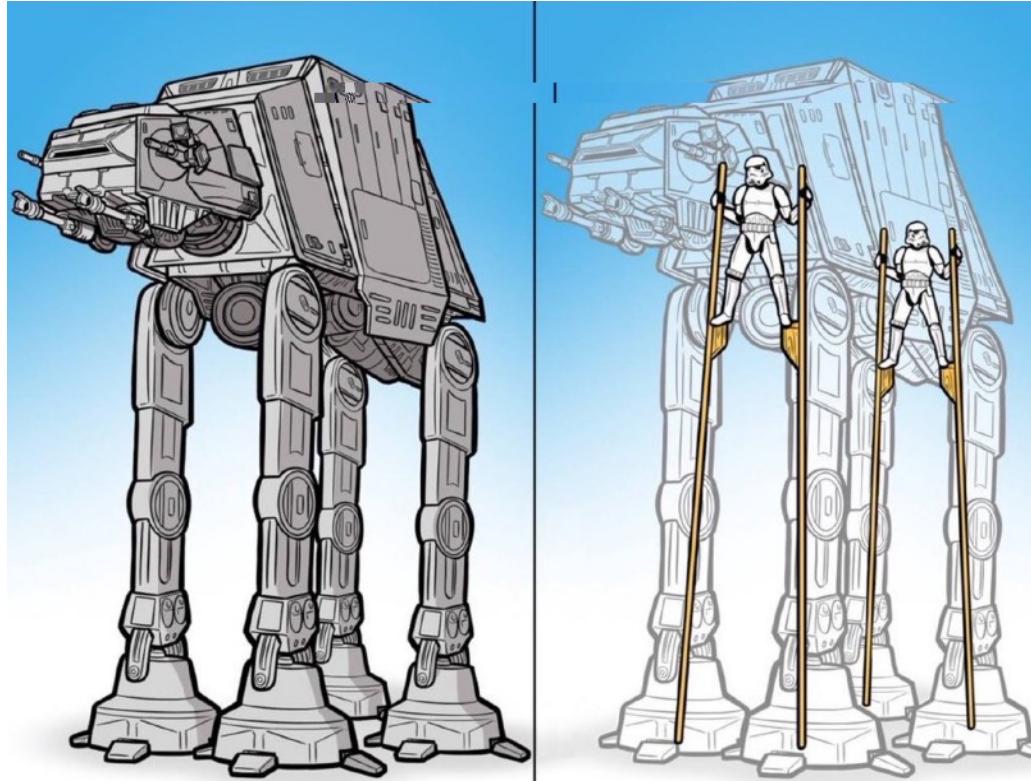
~2000

Today

The Dark Ages

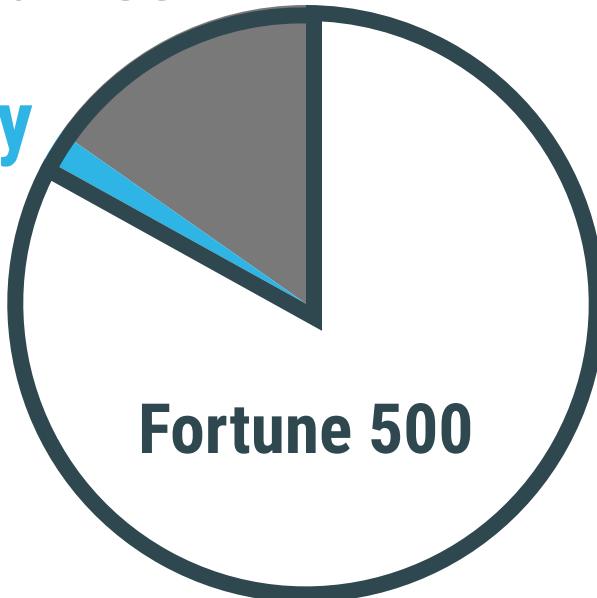




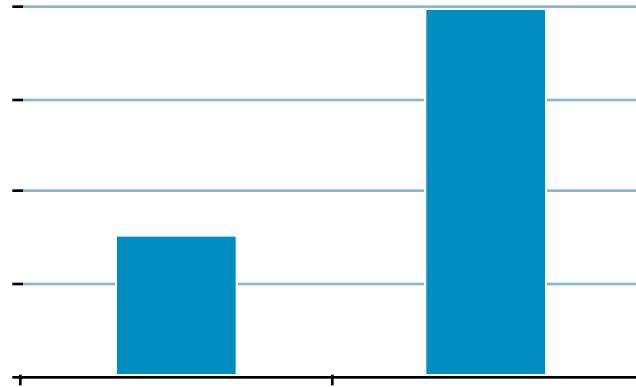


Tech Companies

Good at security



Fortune 500



People
we have People
we need

Building not breaking

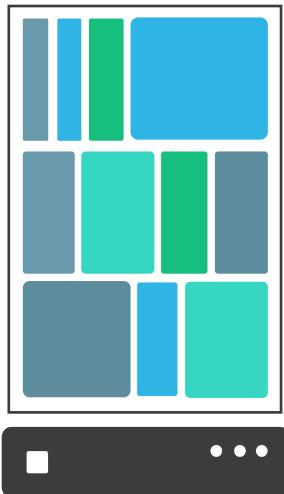
If it doesn't come on by
default, no one will use it



Biased Timeline of the Evolution of Security

~2000

Today

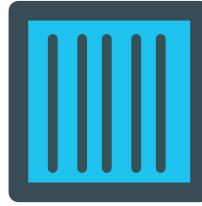


Immutable Infrastructure

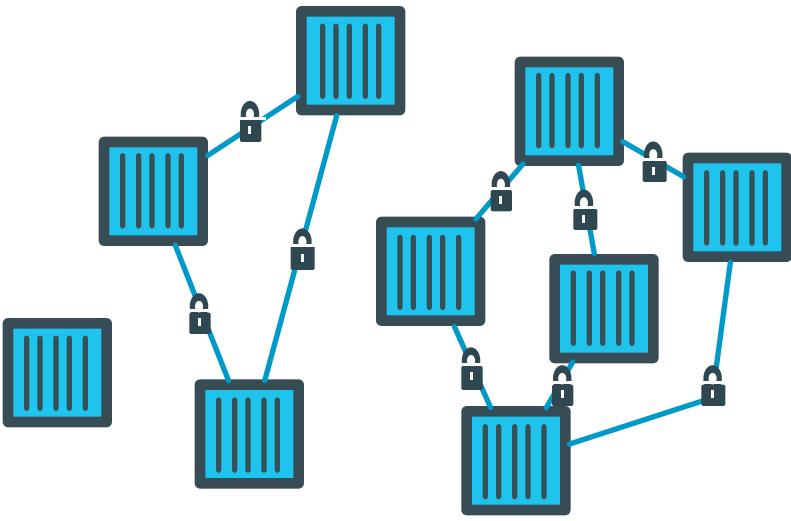
• • •

```
# echo "evil!" > /bin/sh
/bin/sh: can't create /bin/sh: Read-only file system
```

Application Containers



Decentralized Authorization

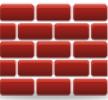


~2000

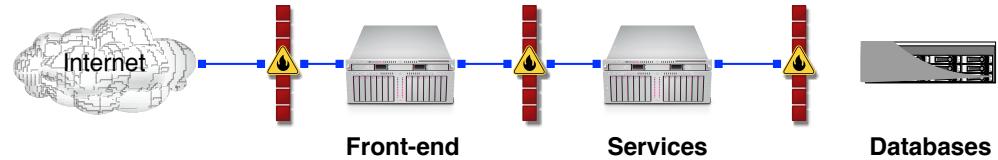
Today



panda



Security
boundary is the
service



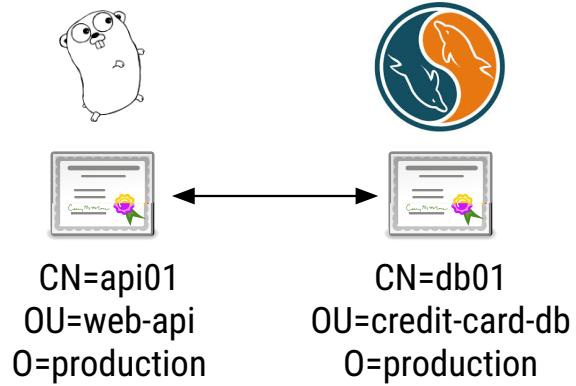
One App Instance, One ID



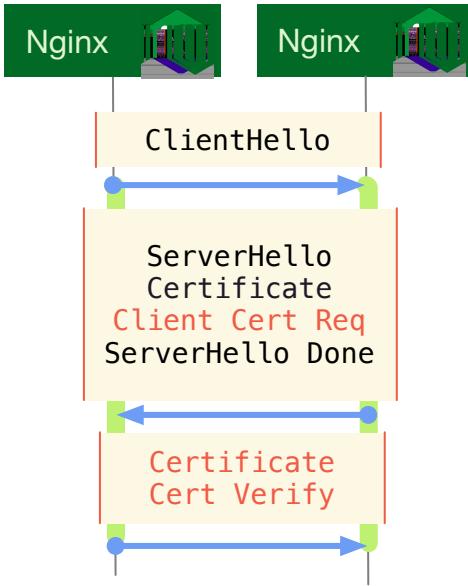
One Node,
One ID



Every service
call is
authorized and
authenticated



mTLS



mTLS - the good

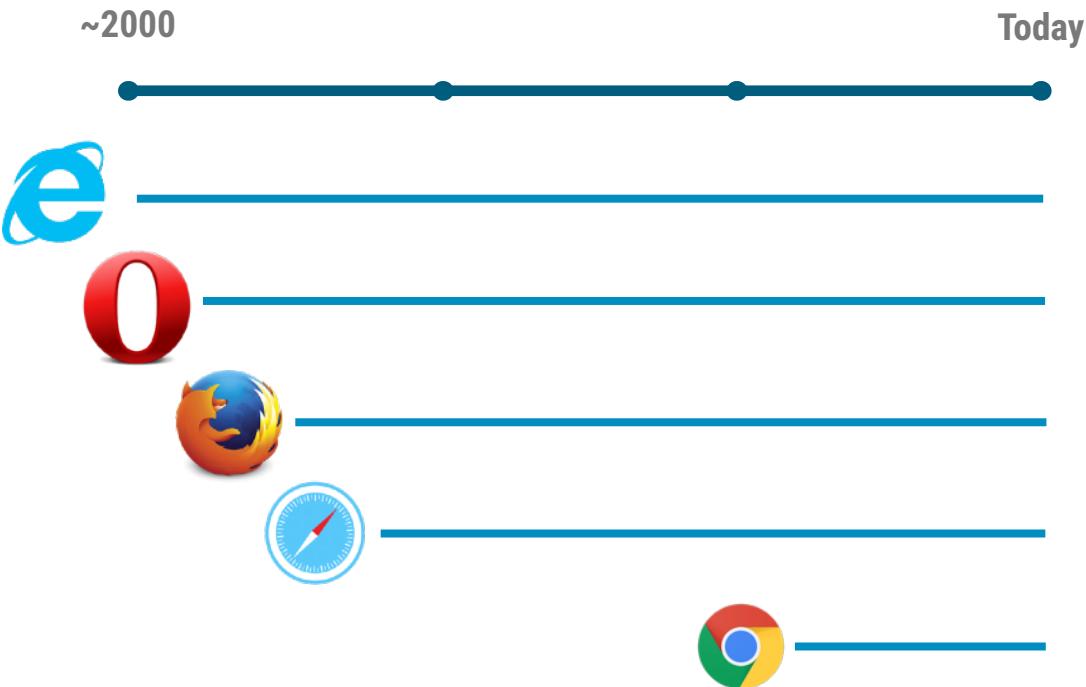
- 1. Key-material stays secret**
- 2. Supported everywhere**
- 3. Authentication and Encryption**

mTLS - the bad

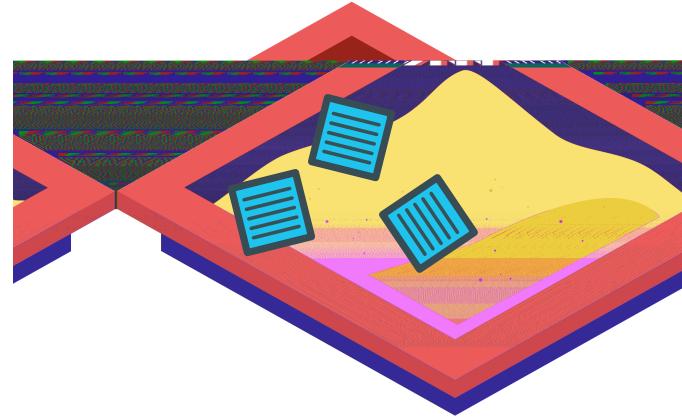
- 1. A LOT of certs**
- 2. Confusing for developers**
- 3. No good revocation story**
- 4. Running a PKI is hard**
- 5. Unforgiving**

Least-privilege resource access

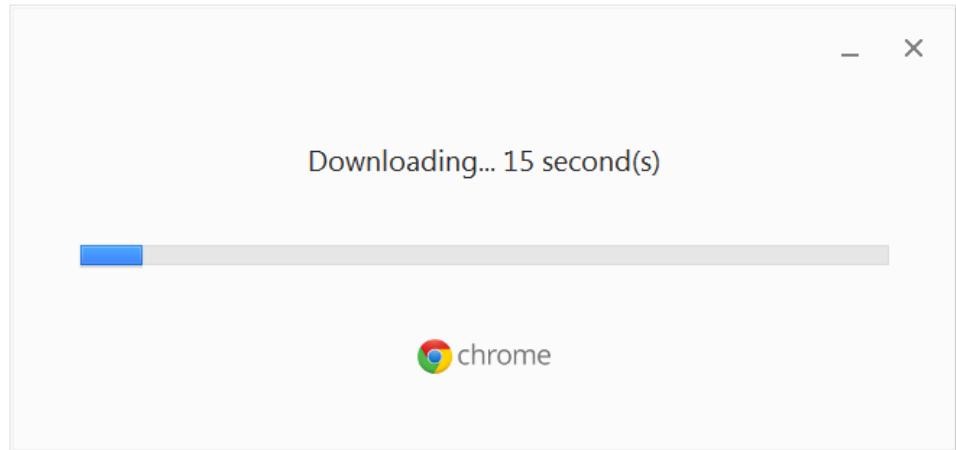
```
[ { "permission":  
    { "method": "GET", "resource": "/user" },  
    "allow": [ "web", "fulfillment", "payments" ] },  
  
  { "permission":  
    { "method": "POST", "resource": "/user" },  
    "allow": [ "signup", "web" ] },  
  
  { "permission":  
    { "method": "DELETE", "resource": "/user/.*" },  
    "allow": [ "web" ]  
  } ]
```



Sandboxing

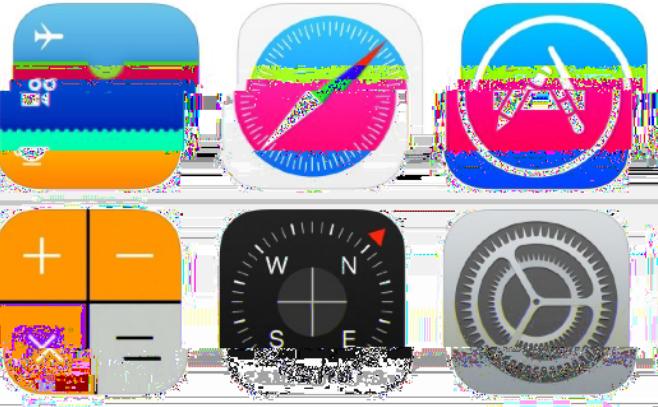


Transparent Updates

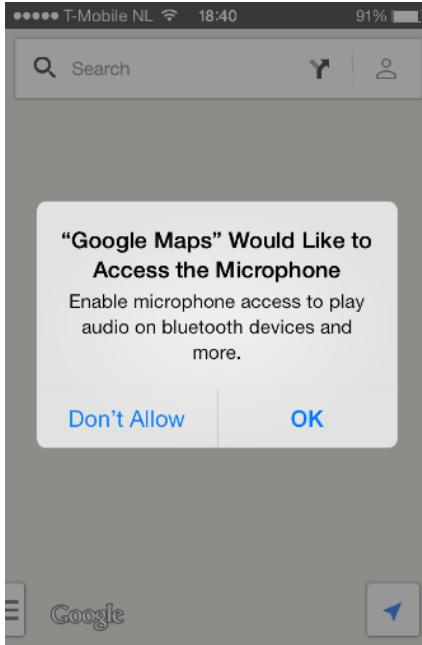




Application Whitelist

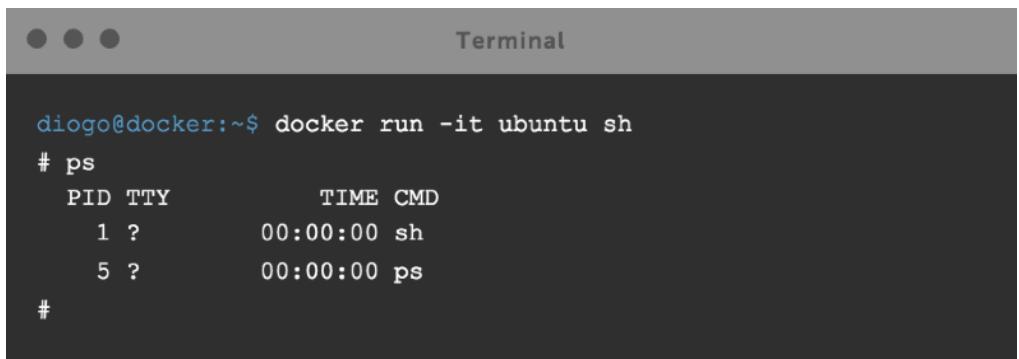


Usable Security



THE UNIVERSITY OF TORONTO LIBRARIES

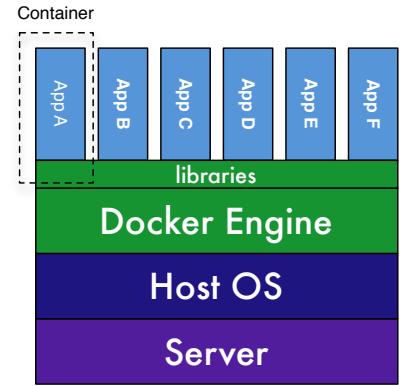
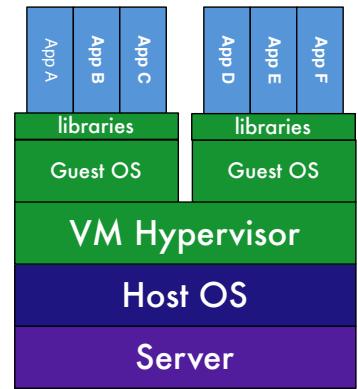
A container is isolated



A screenshot of a Mac OS X terminal window titled "Terminal". The window shows a command-line session on a Ubuntu system within a Docker container. The user has run the command "ps" to list processes. There are two processes listed: one with PID 1, TTY ?, TIME 00:00:00, and CMD sh; and another with PID 5, TTY ?, TIME 00:00:00, and CMD ps. A final "#>" prompt is visible at the bottom.

```
diogo@docker:~$ docker run -it ubuntu sh
# ps
  PID TTY      TIME CMD
    1 ?        00:00:00 sh
    5 ?        00:00:00 ps
#
#
```

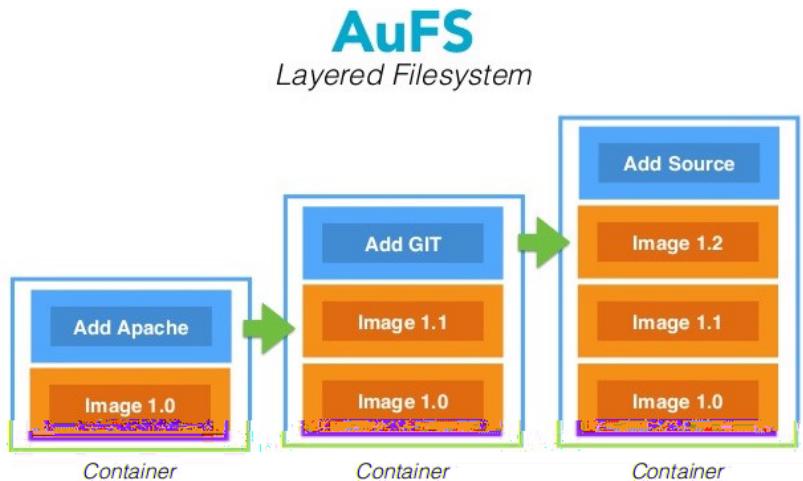
Containers VS VMs



A container is a process

| Process | % time | seconds | usecs/call | calls | errors | syscall |
|---------|--------|----------|------------|-------|--------|--------------|
| 6389 | 62.83 | 0.000747 | 68 | 11 | 6 | wait4 |
| 6476 | 7.91 | 0.000094 | 6 | 17 | | lstat |
| 6477 | 5.05 | 0.000060 | 15 | 4 | | getdents64 |
| 6479 | 3.45 | 0.000041 | 2 | 19 | | open |
| 6480 | 3.20 | 0.000038 | 0 | 106 | | read |
| 6481 | 2.61 | 0.000031 | 31 | 1 | | nanosleep |
| | 2.02 | 0.000024 | 2 | 10 | | mprotect |
| | 1.85 | 0.000022 | 1 | 38 | | rt_sigaction |
| | 1.60 | 0.000019 | 0 | 43 | | ioctl |
| | 1.43 | 0.000017 | 1 | 20 | | close |
| | 1.26 | 0.000015 | 0 | 83 | | writev |

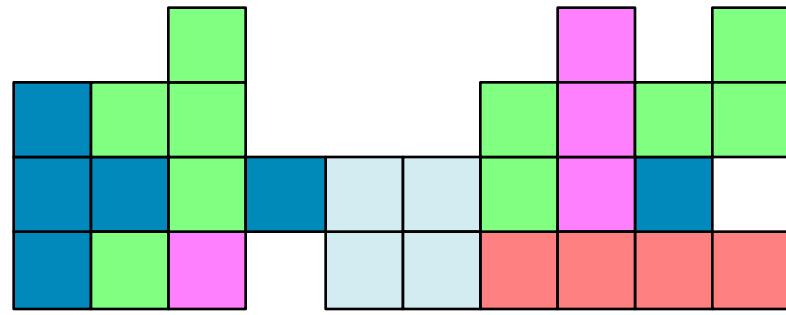
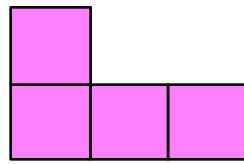
Copy-on-write FS

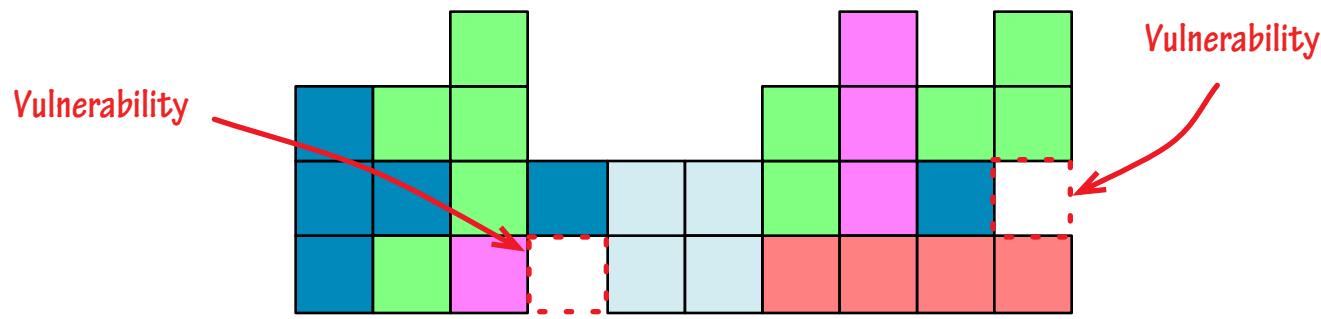
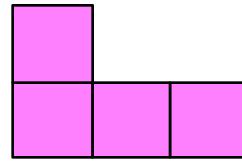


Key Security advantages

- 1. Additional layer of isolation**
- 2. Easier patch management, faster deployments, reduced risk**
- 3. Better visibility into application behavior**

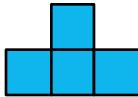
```
...  
$ docker stack deploy --stack-file=socket-stack.yml my_app  
Creating service my_app_frontend  
Creating service my_app_backend  
Creating service my_app_db
```



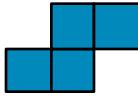


Security Tetrominos

runC



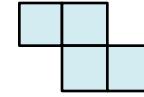
containerD



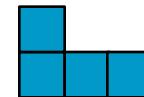
Docker



infraKit



linuxKit



Notary

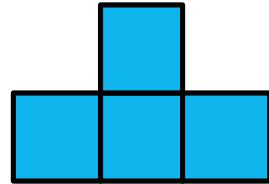


swarmKit



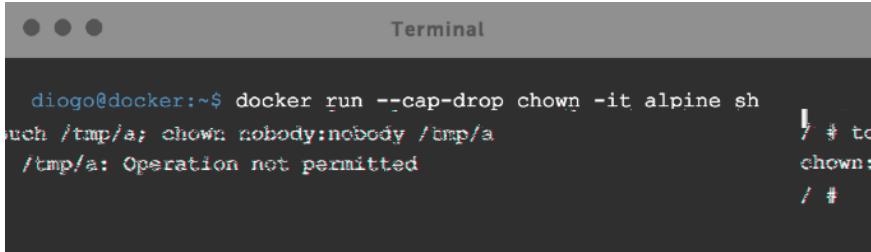
runC

Lightweight universal container runtime



runC

Capabilities



A screenshot of a terminal window titled "Terminal". The window shows a command being run: `diogo@docker:~$ docker run --cap-drop chown -it alpine sh`. The user attempts to change ownership of a file using the command `chown /tmp/a; chown nobody:nobody /tmp/a`. The output shows that the first part of the command succeeds, but the second part fails with the error message `/tmp/a: Operation not permitted`. The terminal prompt is visible at the bottom right.

```
diogo@docker:~$ docker run --cap-drop chown -it alpine sh
chown /tmp/a; chown nobody:nobody /tmp/a
/tmp/a: Operation not permitted
/ #
```

runC

- Namespace Isolation
- Cgroups

Namespaces



Cgroups





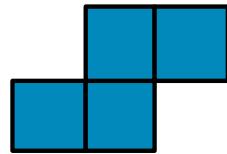
runC

Seccomp-bpf

```
{  
    "defaultAction": "SCMP_ACT_ERRNO",  
    "architectures": [  
        "SCMP_ARCH_X86_64",  
    ],  
    "syscalls": [  
        {  
            "name": "accept",  
            "action": "SCMP_ACT_ALLOW",  
            "args": []  
        },  
        {  
            "name": "accept4",  
            "action": "SCMP_ACT_ALLOW",  
            "args": []  
        },  
        ...  
    ]  
}
```

containerD

Container runtime supervisor



containerD

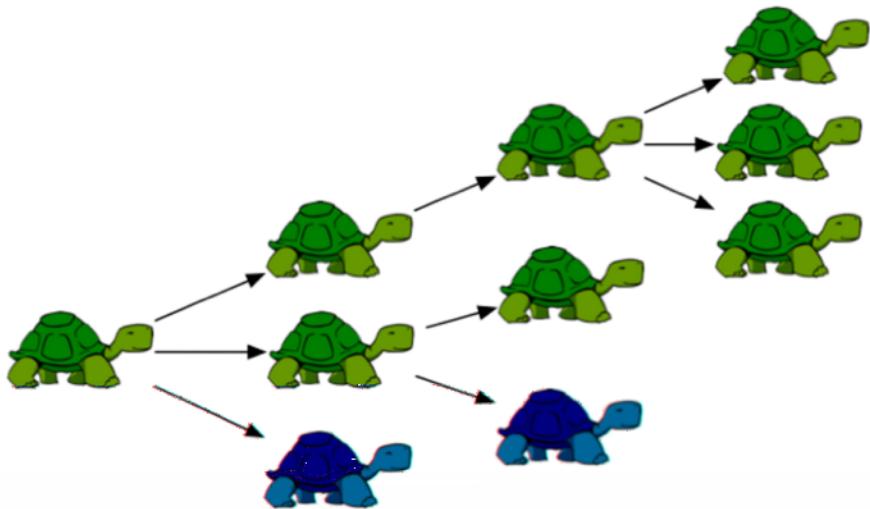
Content
addressable
images

H(blob) ==
3624bbc5e6074c



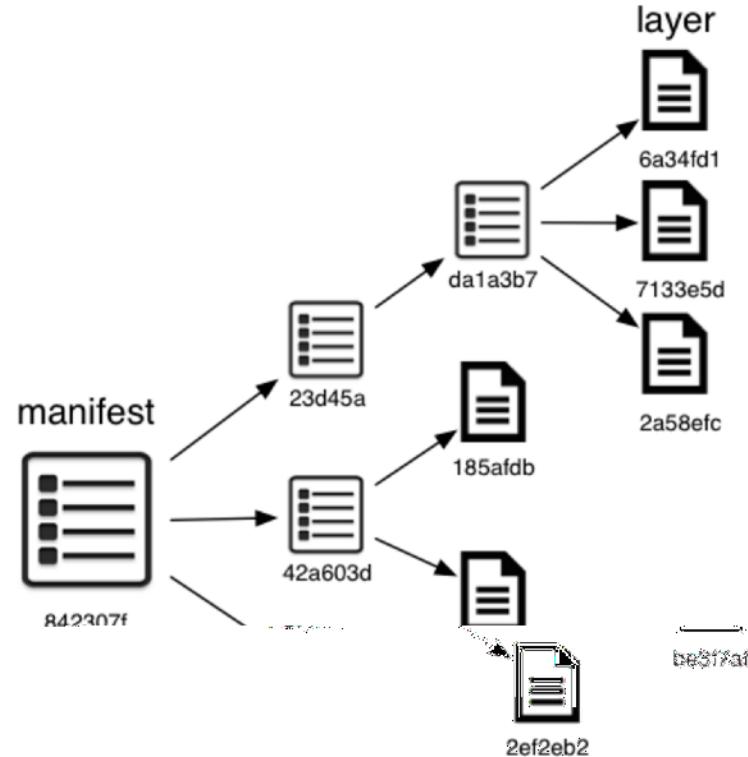
containerD

Turtles all the way
down



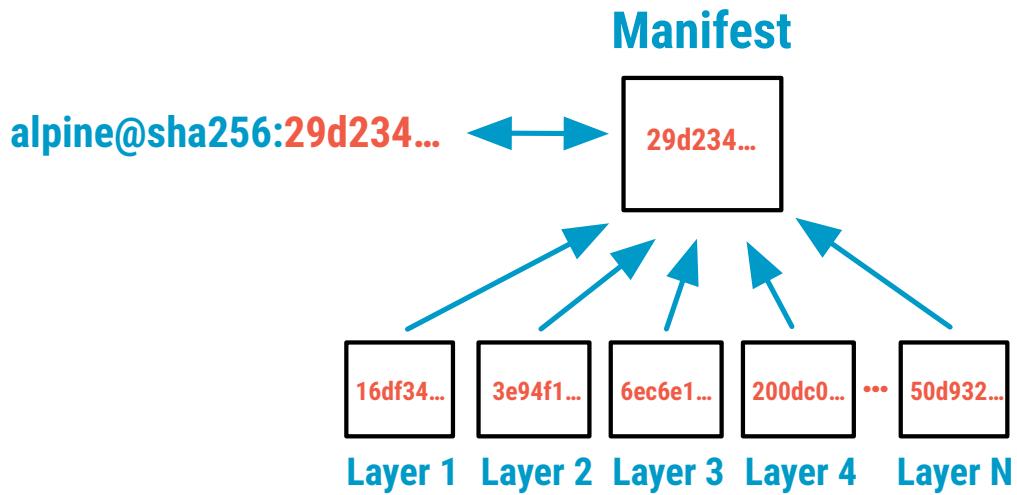
containerD

Images are Merkle DAGs of layers



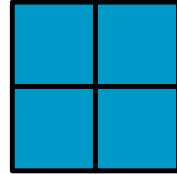
containerD

Content Addressable
Image Pulls



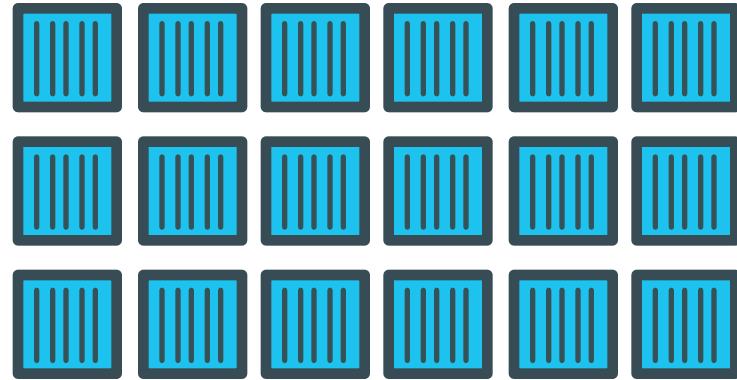
Docker

**Secure-by-default software
container platform**



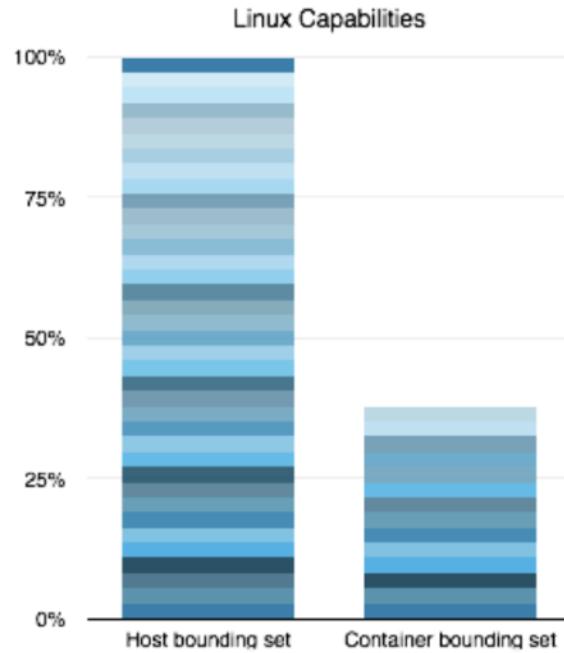
Docker

- SELinux & AppArmor
- Capability Whitelist
- Syscall Whitelist
- Copy-on-write
- No device access



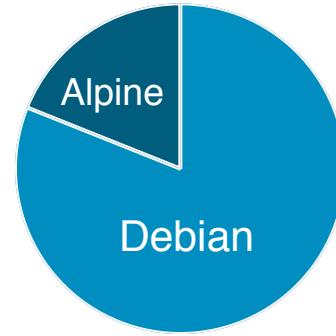
Docker

Less than half the default capabilities



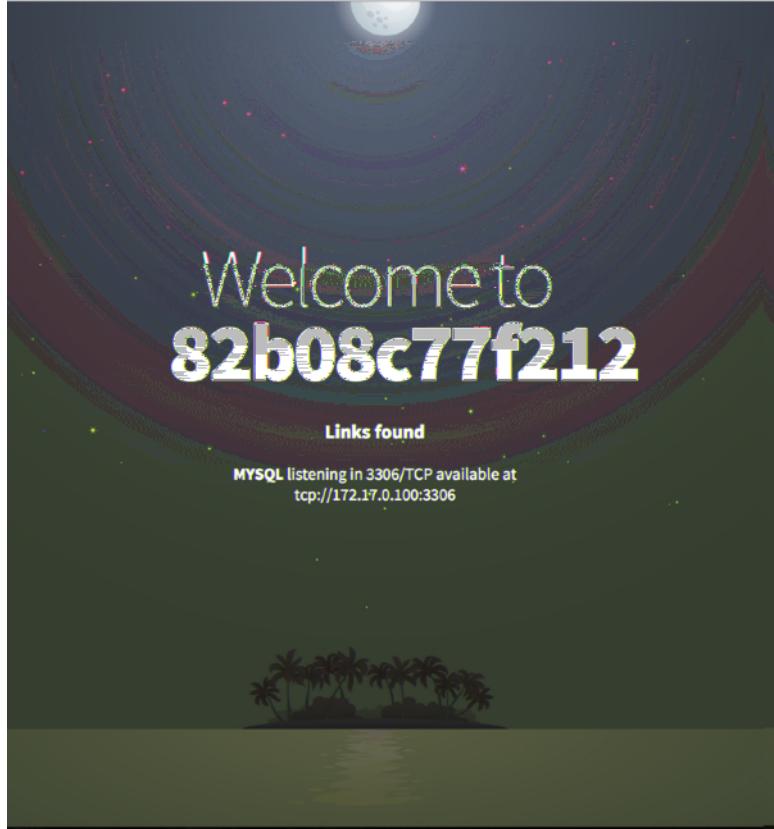
Docker

Minimal Distros



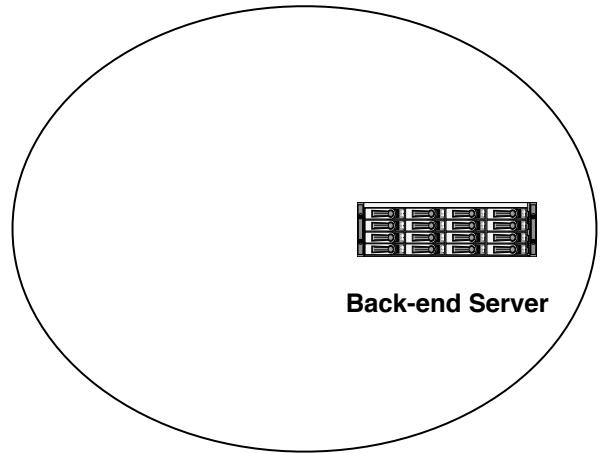
Demo

Our new website



Demo

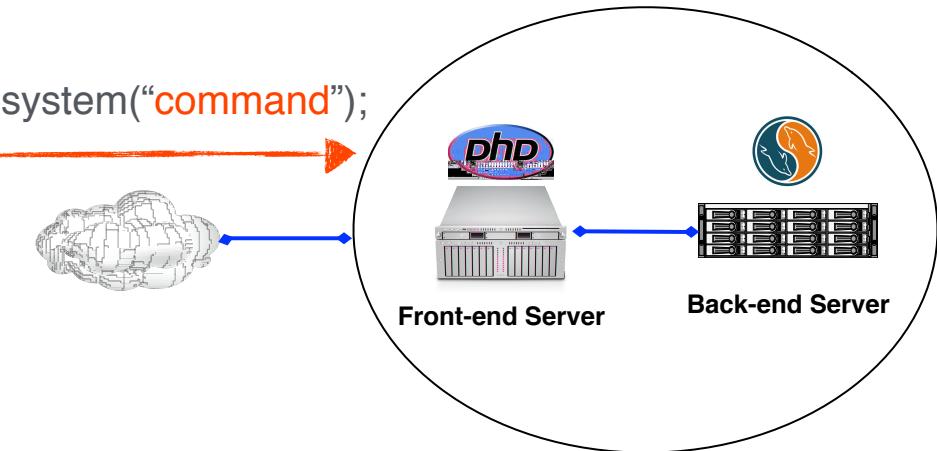
Our new architecture



Demo

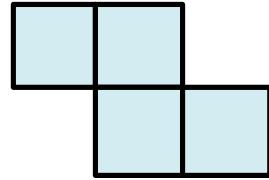
Our new backdoor

```
?shell=system("command");
```



infraKit

Infrastructure independent machine management



infraKit

Platform Agnostic



infraKit

Declarative Updates

```
# infrakit group commit moby.json
Performing a rolling update on 5 instances
```

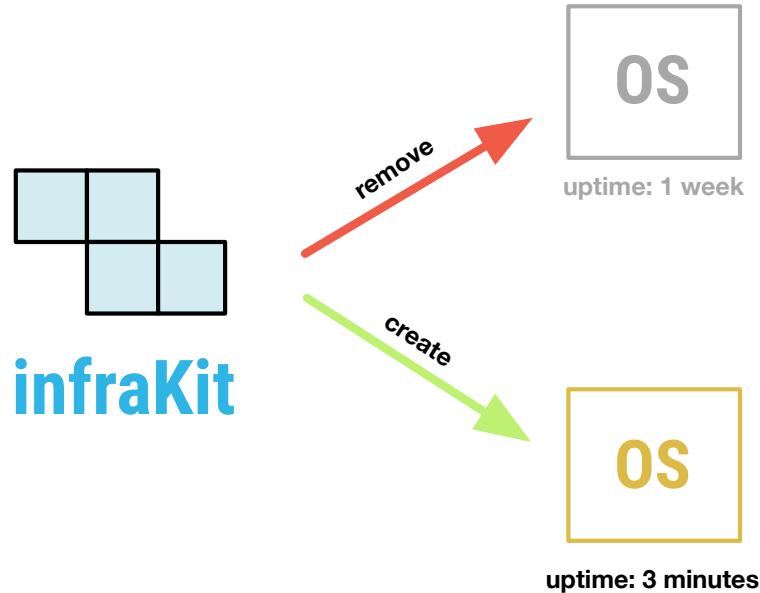
infraKit

Reverse Uptime

10:00

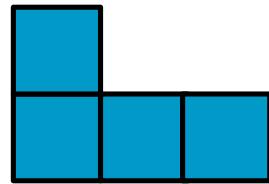
infraKit

Rolling Deploys



linuxKit

A **secure OS builder for your containers**





Immutable Linux OS builder

```
kernel:  
    image: "linuxkit/kernel:4.9.x"  
    cmdline: "console=ttyS0 page_poison=1"  
init:  
    - linuxkit/init:1b8a7e394d...  
onboot:  
    - name: dhcpcd  
        image: "linuxkit/dhcpcd:7d2b8aaaf..."  
        command: ["/sbin/dhcpcd", "--nobackground"]  
trust:  
    org:  
        - linuxkit
```



linuxKit

Minimal Base



linuxKit

Type-safe
System
Daemons



linuxKit

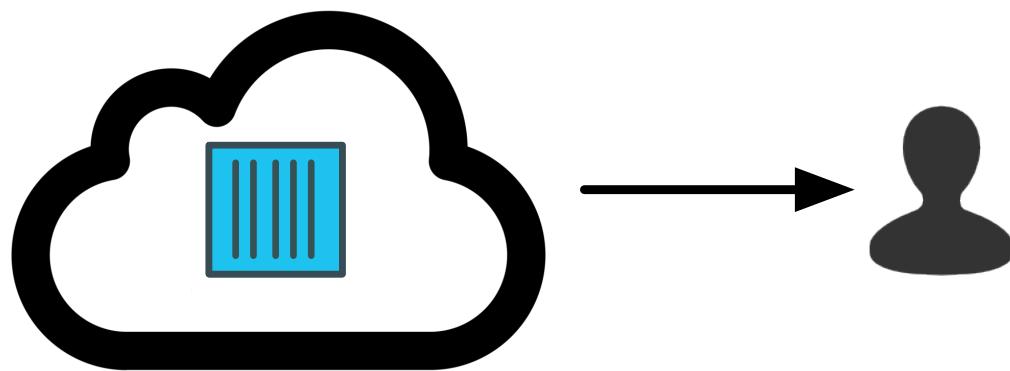
Hardening the kernel

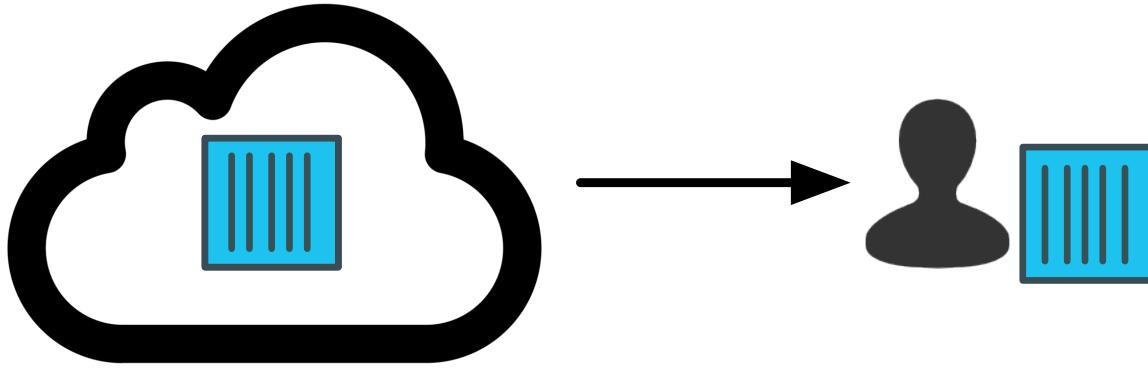


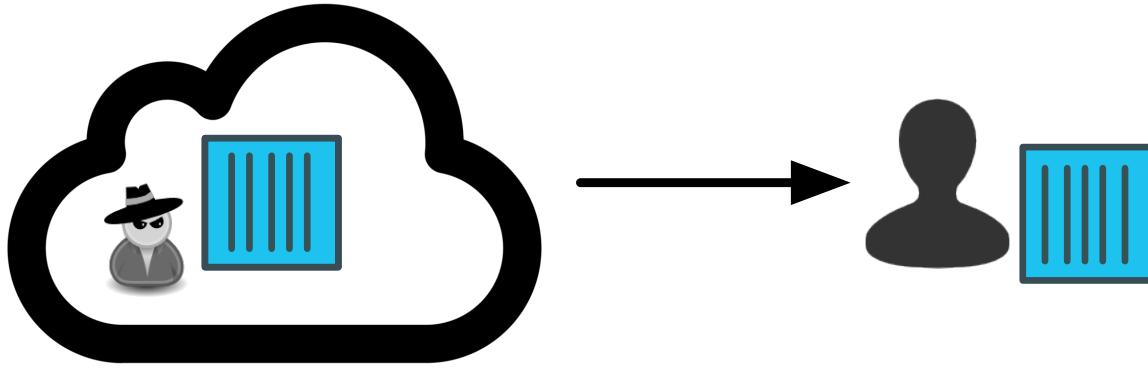
Notary

Trusted software delivery

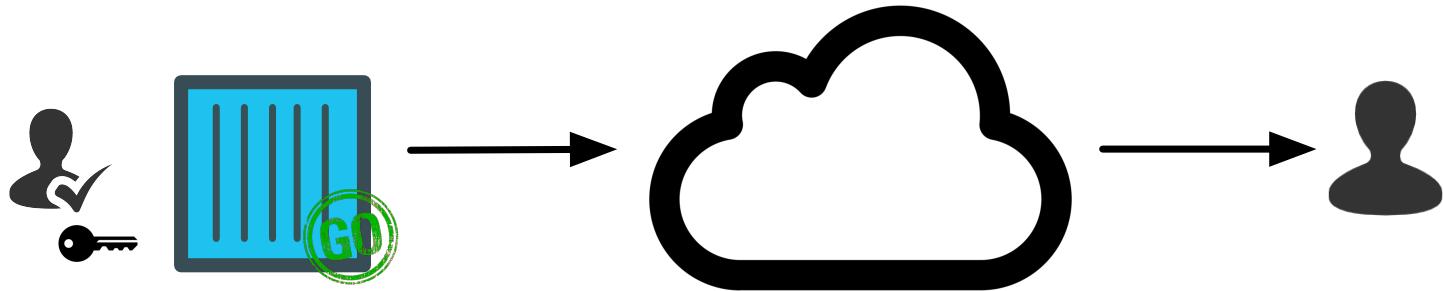


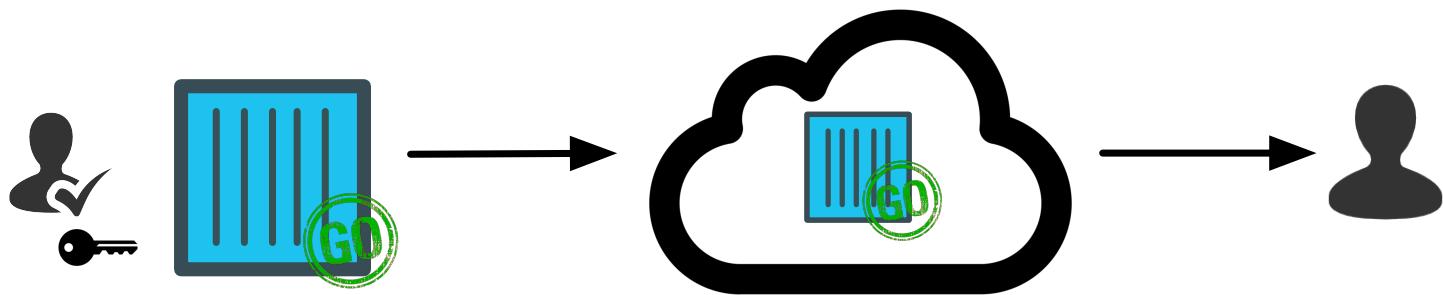


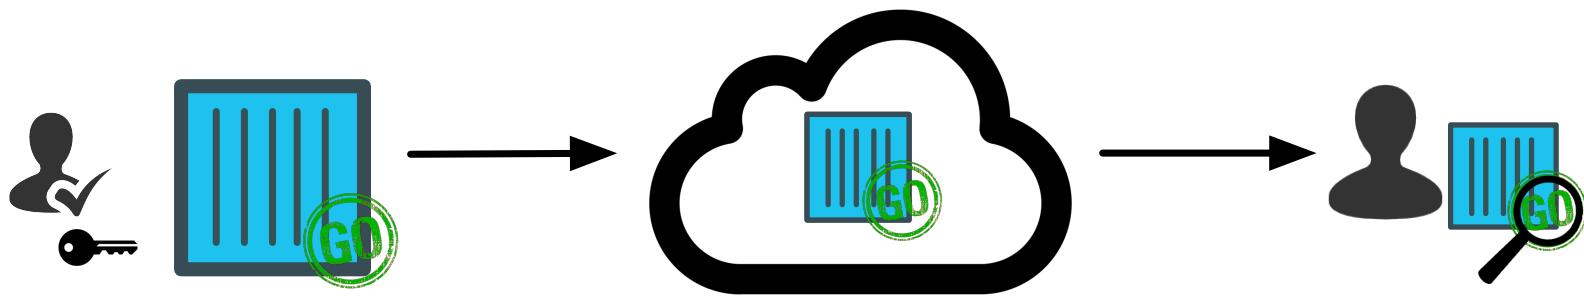


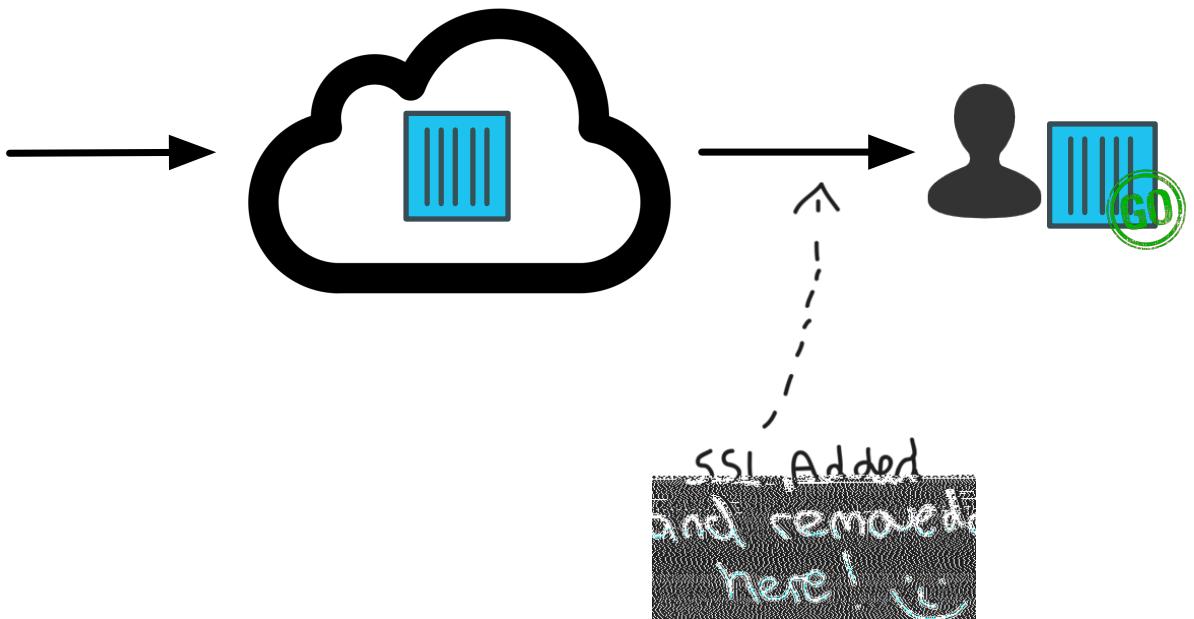


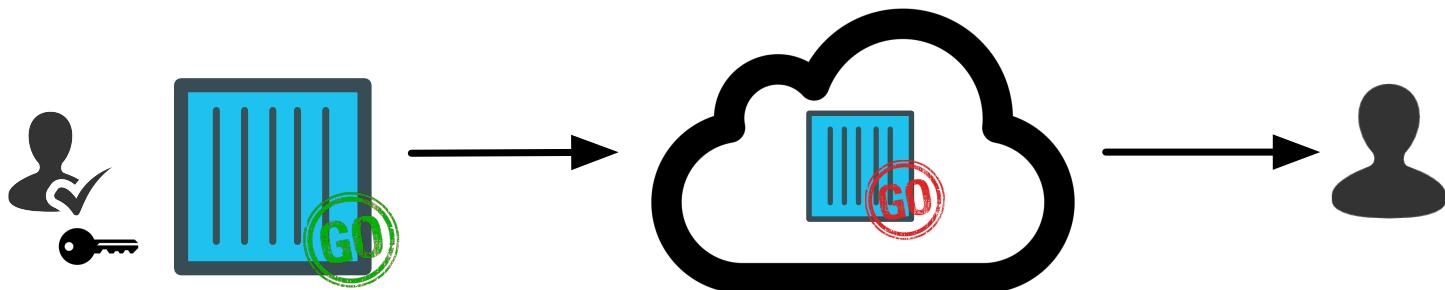
What about GPG?

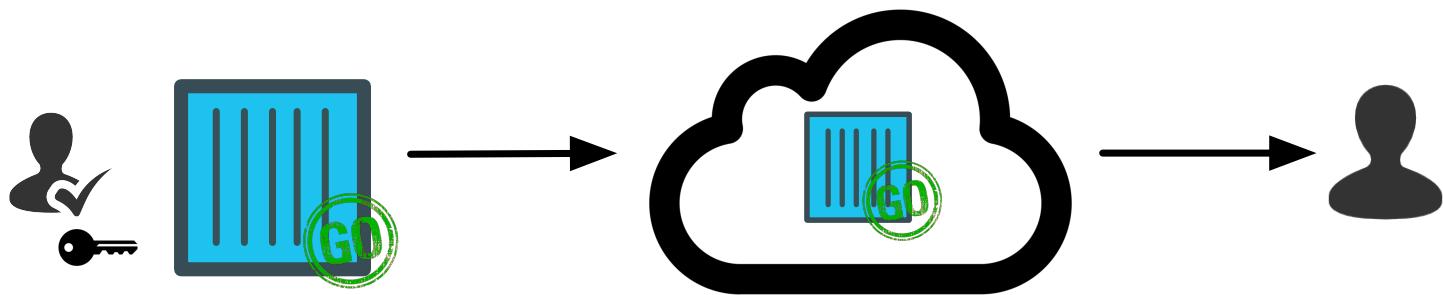


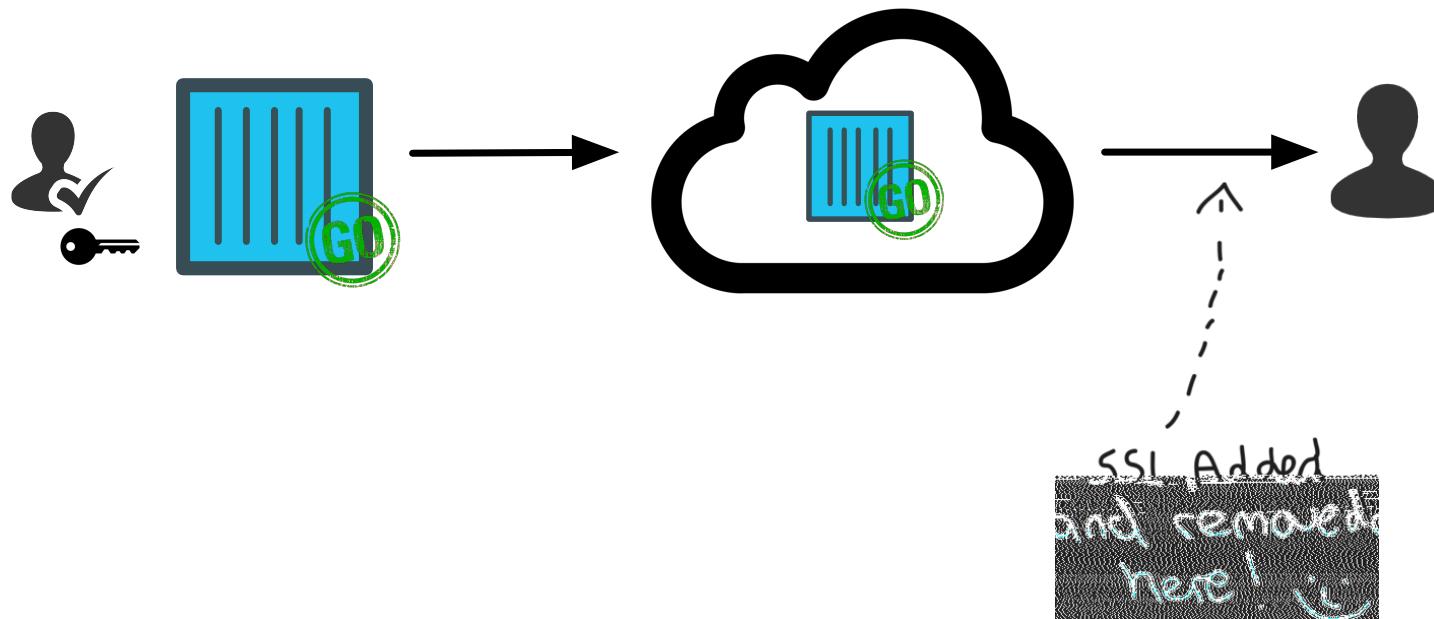


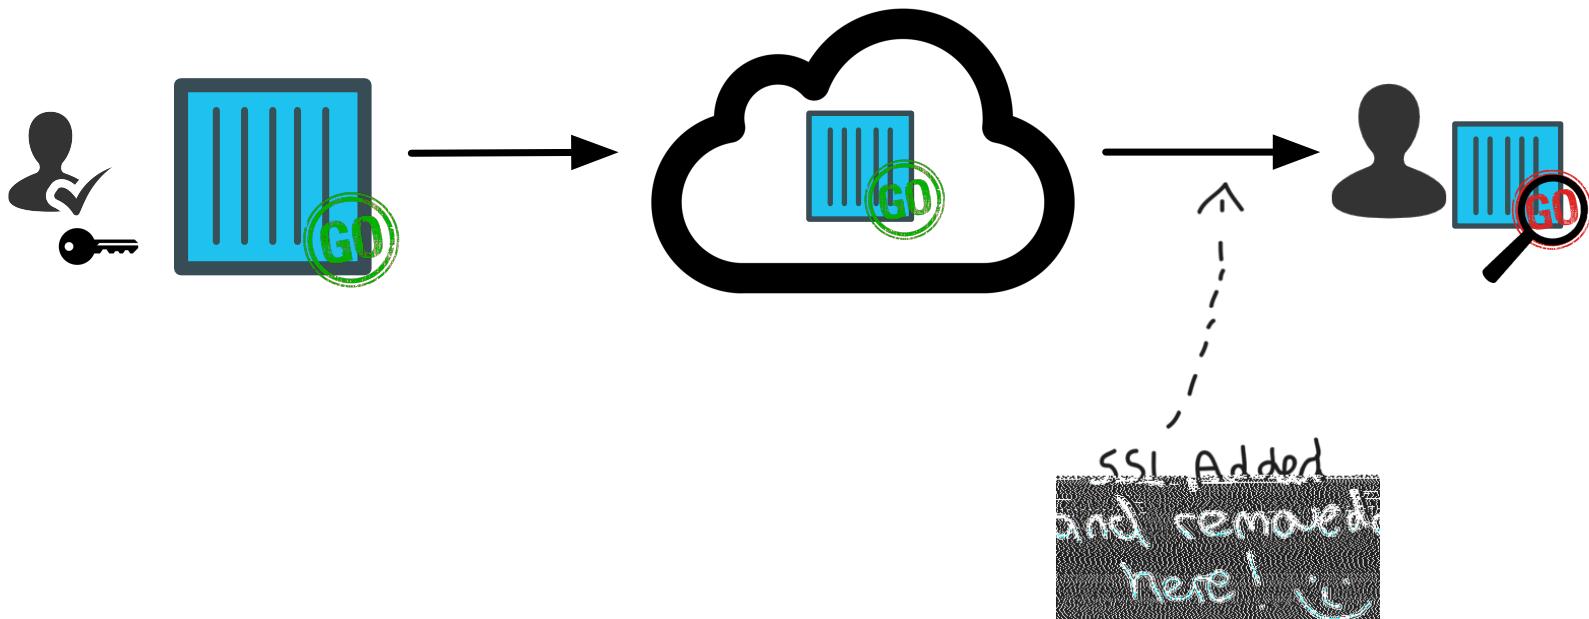












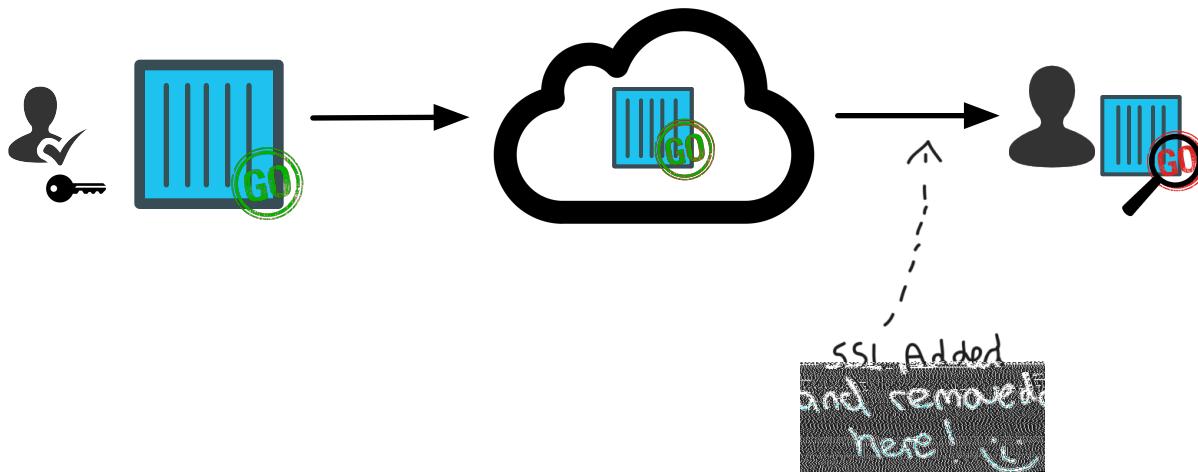


“A software update system is secure if it can be sure that it knows about the latest available updates in a timely manner, any files it downloads are the correct files, and no harm results from checking or downloading files.”

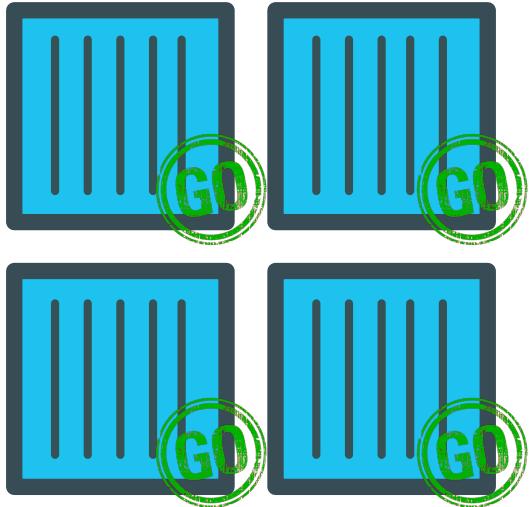
- The Update Framework

- 1. Freshness**
- 2. Signed Collections**
- 3. Key Hierarchy**
- 4. Transparent Key Rotation**
- 5. Threshold Signing**

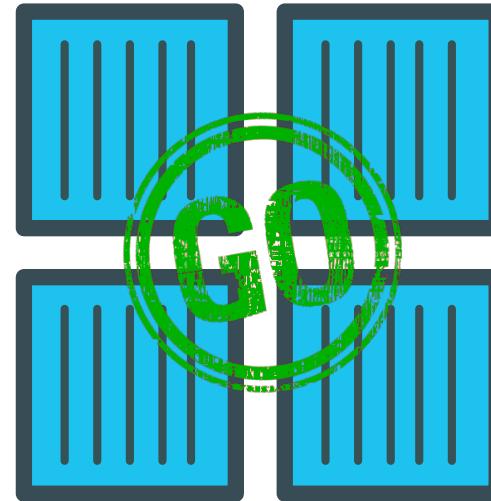
Freshness



Signed Collections

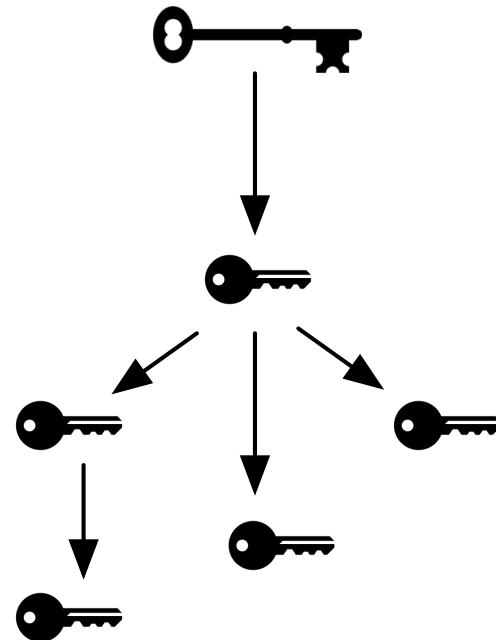


Signed Packages

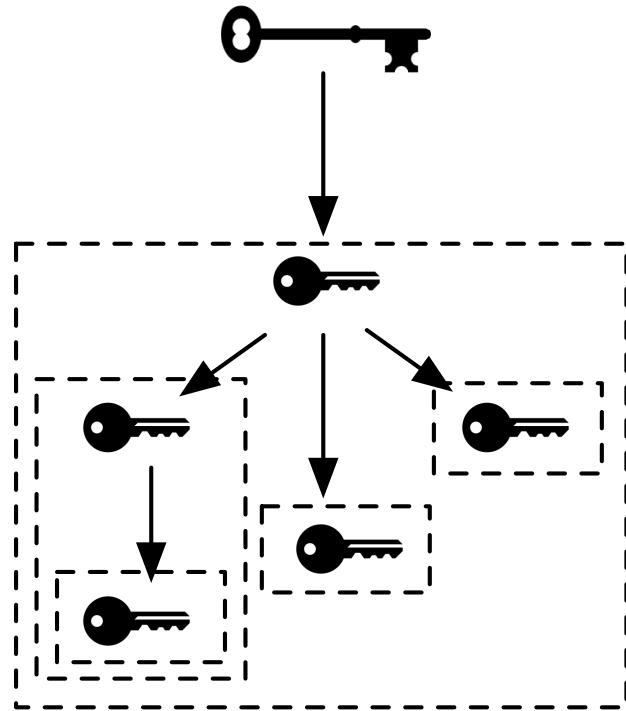


Signed Collection

Key Hierarchy



Key Hierarchy

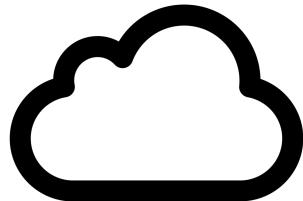


Key Hierarchy

Short
Expiry



Less
Sensitive



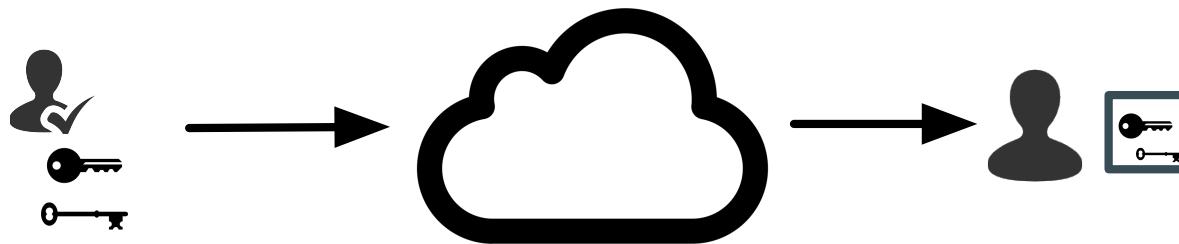
Long
Expiry



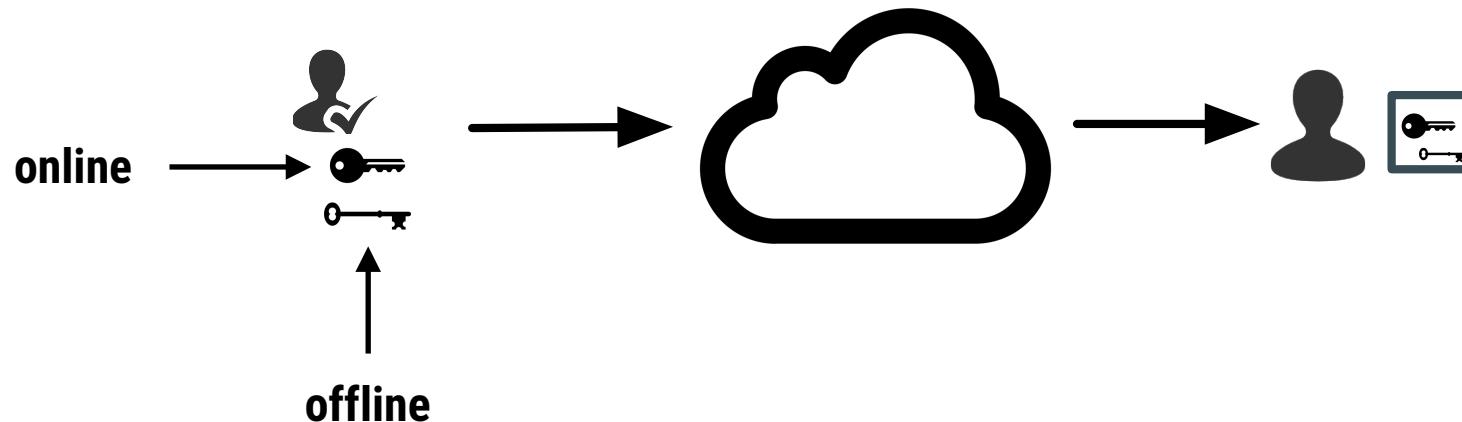
More
Sensitive



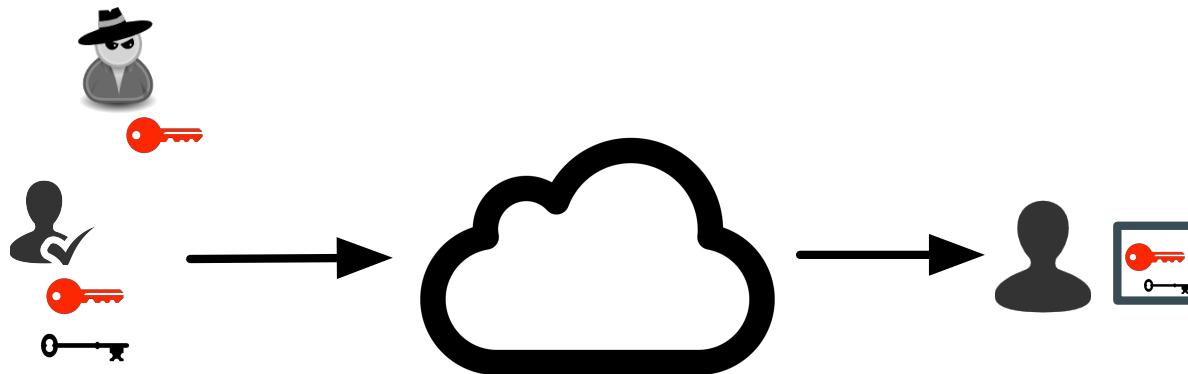
Transparent Key Rotation



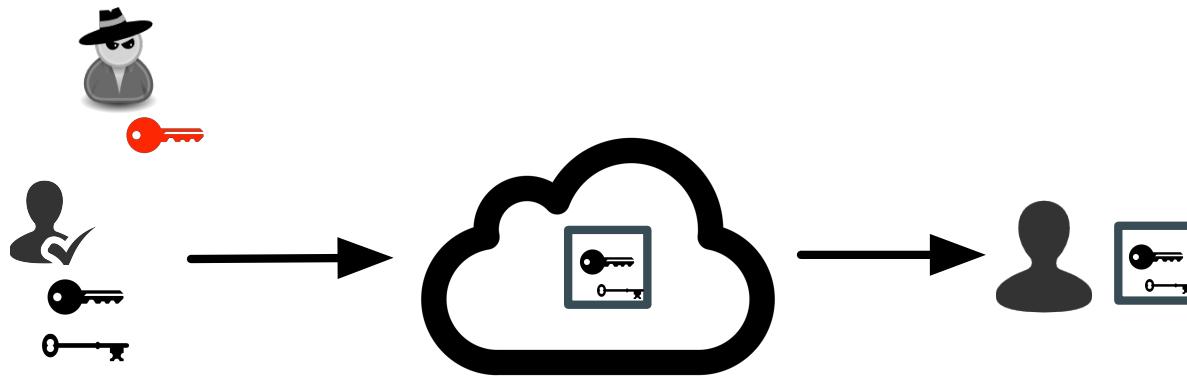
Transparent Key Rotation



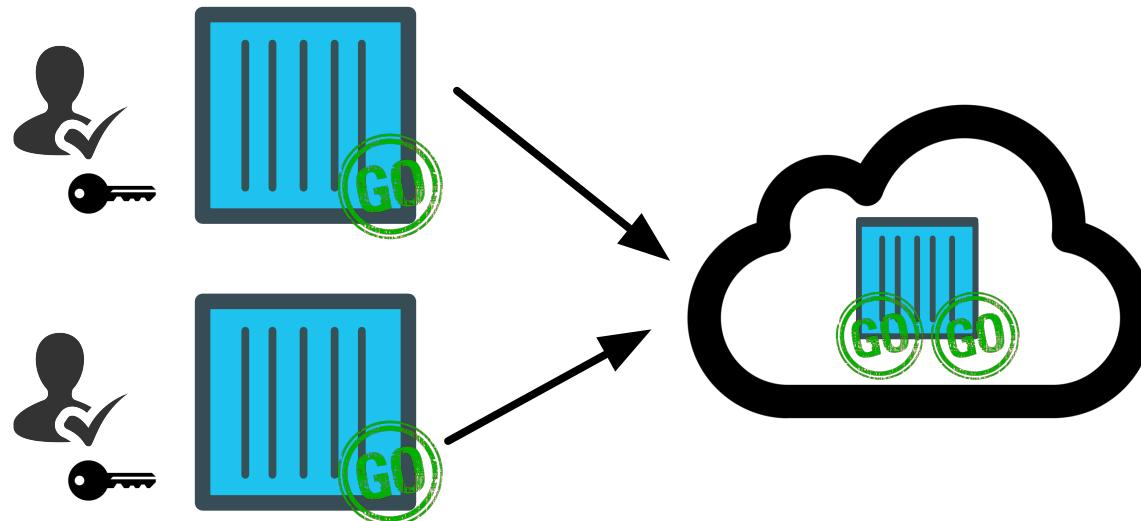
Transparent Key Rotation



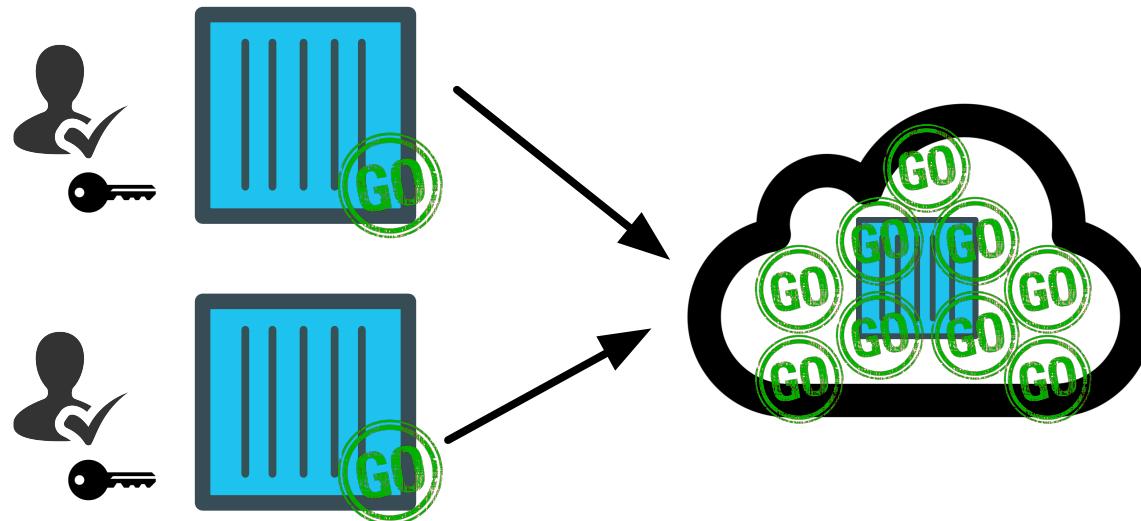
Transparent Key Rotation



Threshold Signing



Threshold Signing



The Update Framework (**TUF**)

theupdateframework.com

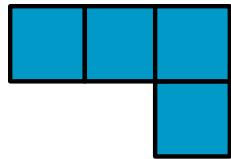
Notary

Cryptographic
Name Resolution



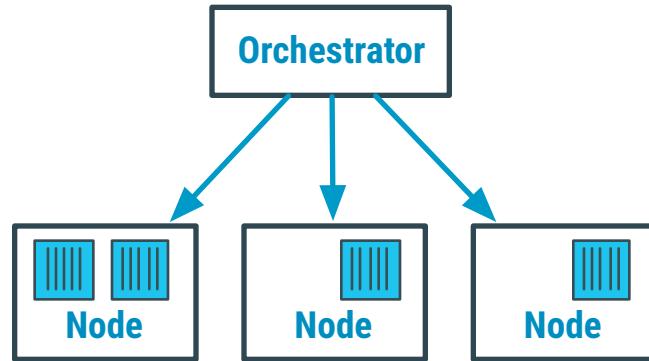
swarmKit

Least-privilege container orchestrator



swarmKit

Multiple
nodes



swarmKit

Secure Node Introduction



swarmKit

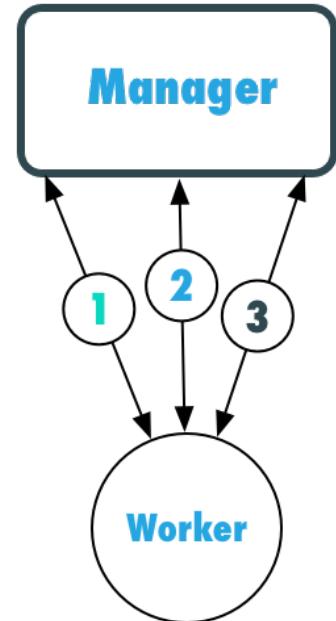
Cryptographic Node Identity

```
$ openssl x509 -in  
/var/lib/docker/swarm/certificates/swarm-node.crt  
-text  
-----  
Swarm ID  
-----  
Cert  
...  
Issuer: CN=swarm-ca  
Validity  
Not Before: Mar  9 15:21:00 2017 GMT  
Not After : Jun  7 16:21:00 2017 GMT  
Subject: O=lgz5xj1eqg... OU=swarm-manager, CN=oFCm6bdy...  
...  
X509v3 Subject Alternative Name:  
DNS:swarm-manager, DNS:grafana, DNS:swarm-13  
...  
-----BEGIN CERTIFICATE-----  
MIICNDCCAdugAwIBAgIUCoRaj23j4h5  
...
```

swarmKit

Bootstrap

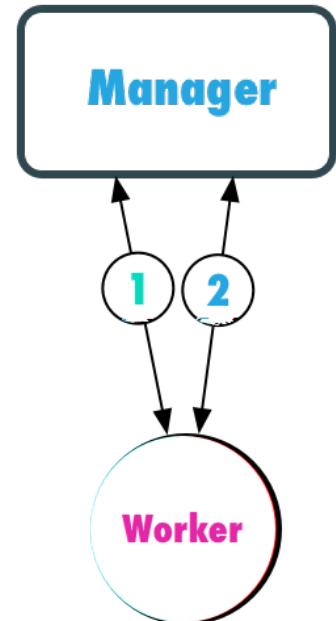
1. Retrieve and validate Root CA Public key material.
2. Submit new CSR along with secret token.
3. Retrieve the signed certificate.



swarmKit

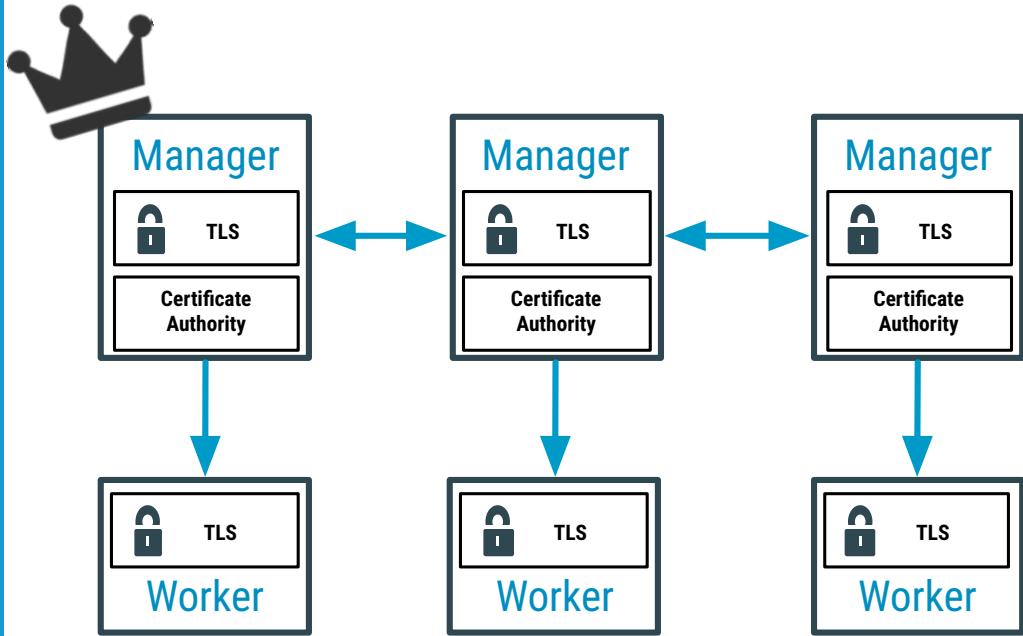
Automatic Certificate Rotation

1. Submit new CSR using old key-pair.
2. Retrieve the new signed certificate.



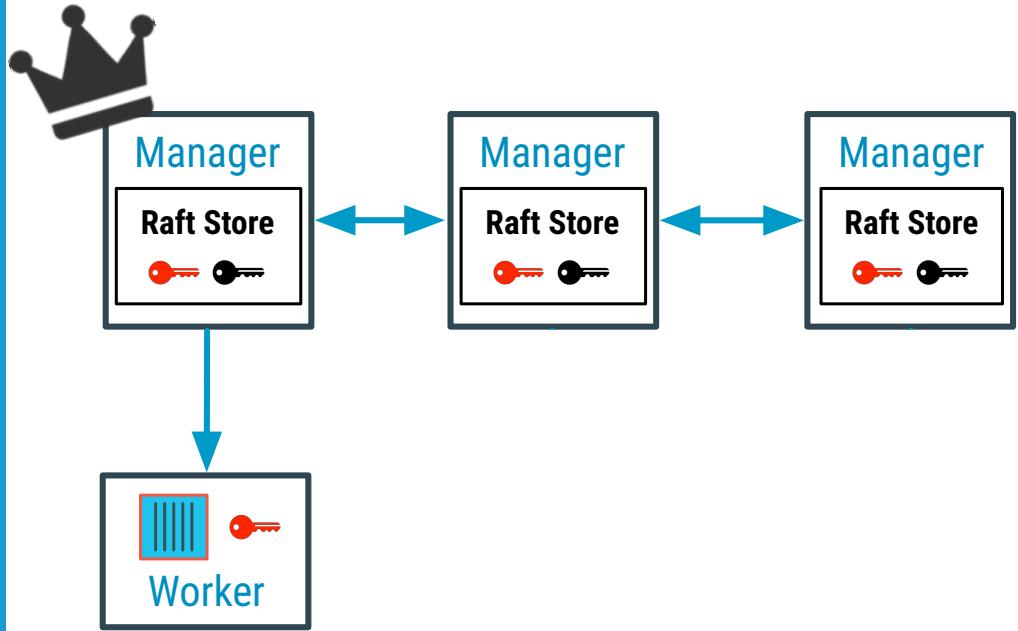
swarmKit

MTLS Between
All Nodes



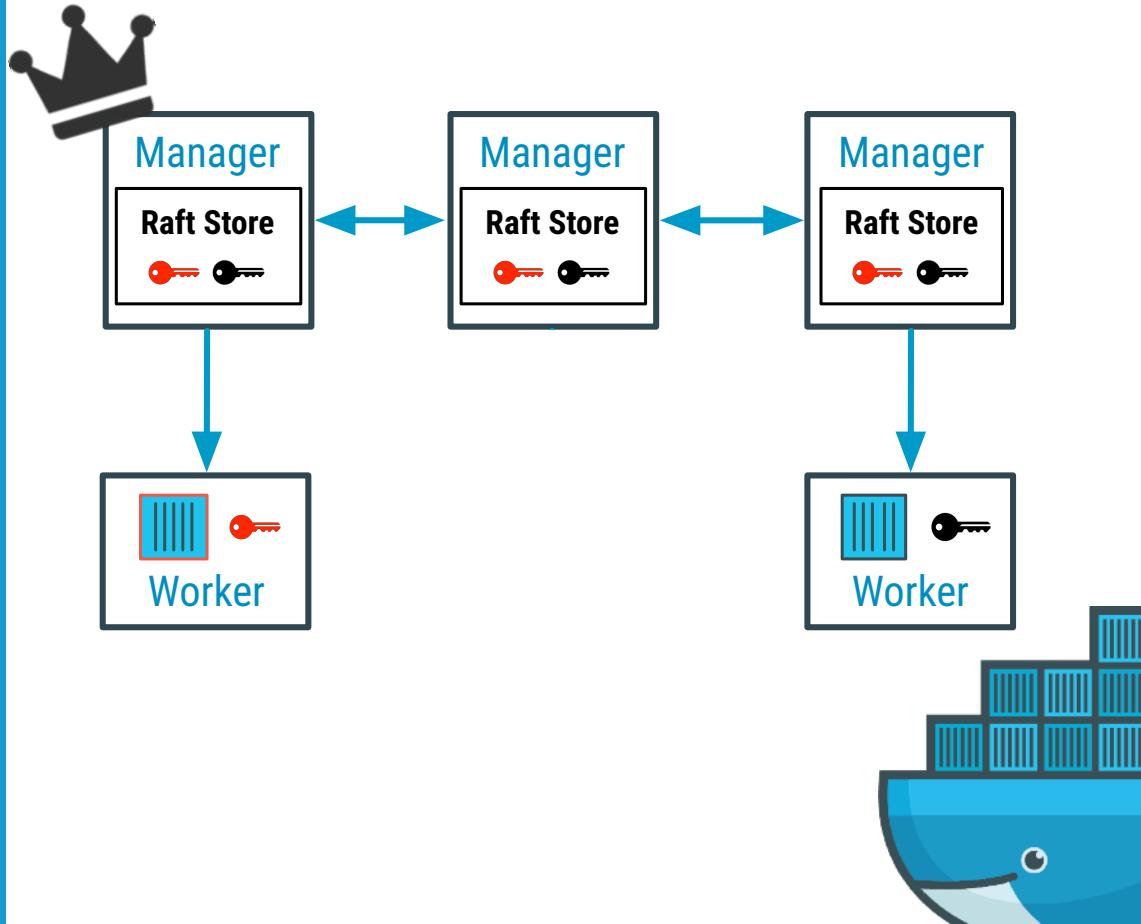
swarmKit

Least-privilege
Secret
Distribution



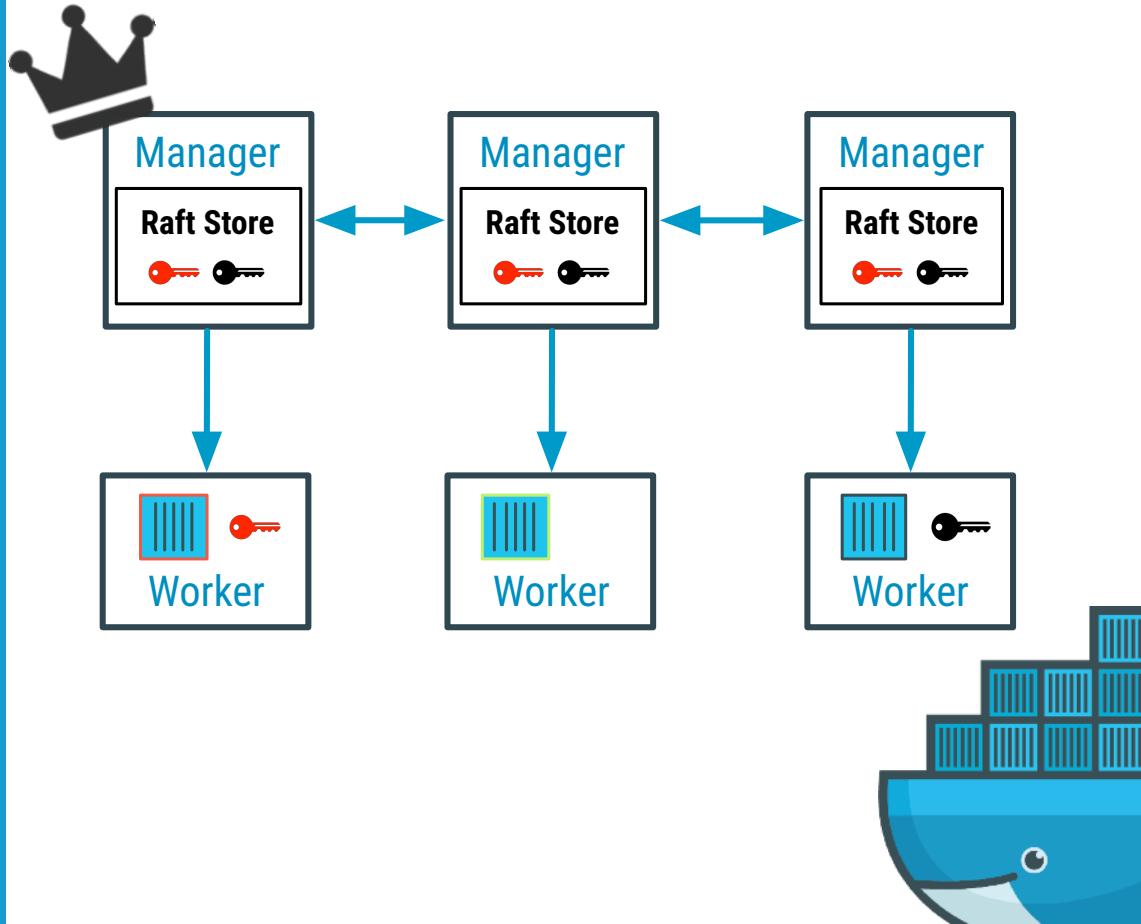
swarmKit

Least-privilege
Secret
Distribution



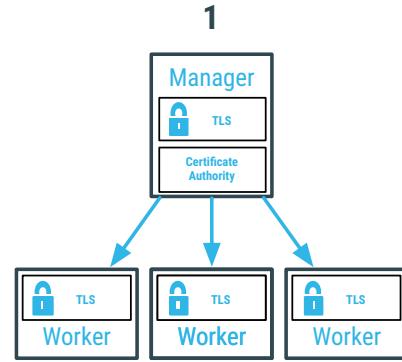
swarmKit

Least-privilege
Secret
Distribution



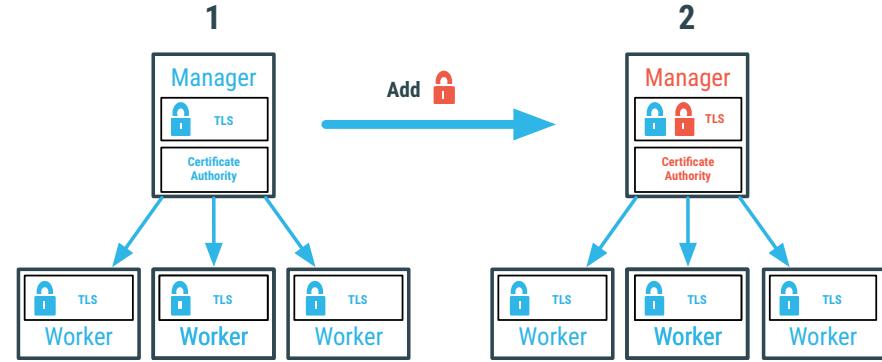
swarmKit

Transparent Root Rotation



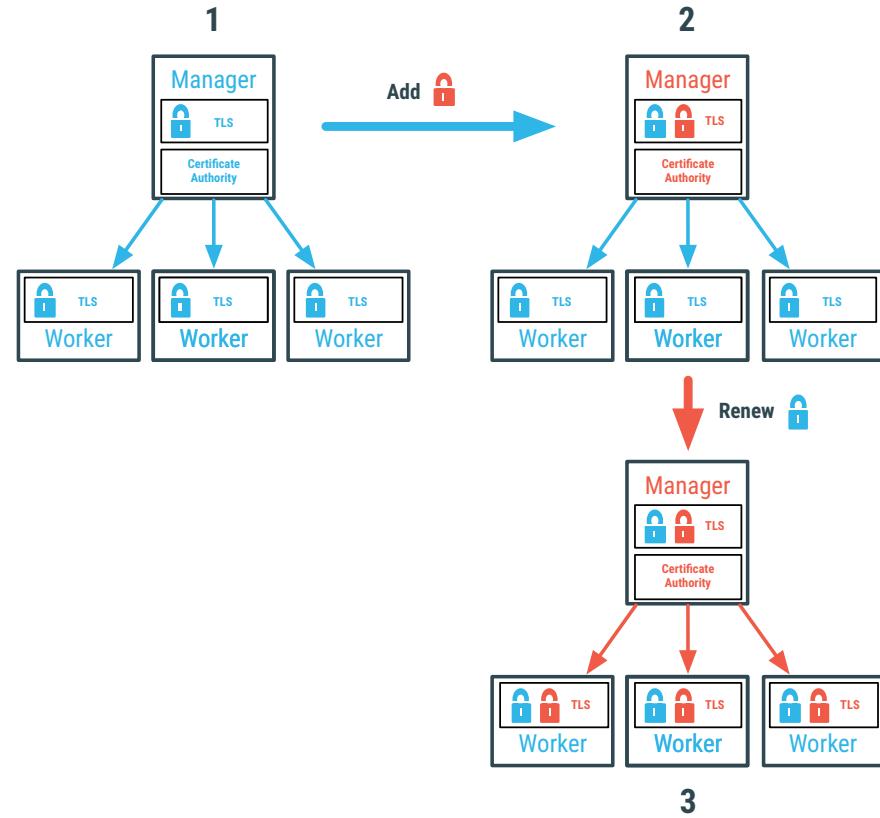
swarmKit

Transparent Root Rotation



swarmKit

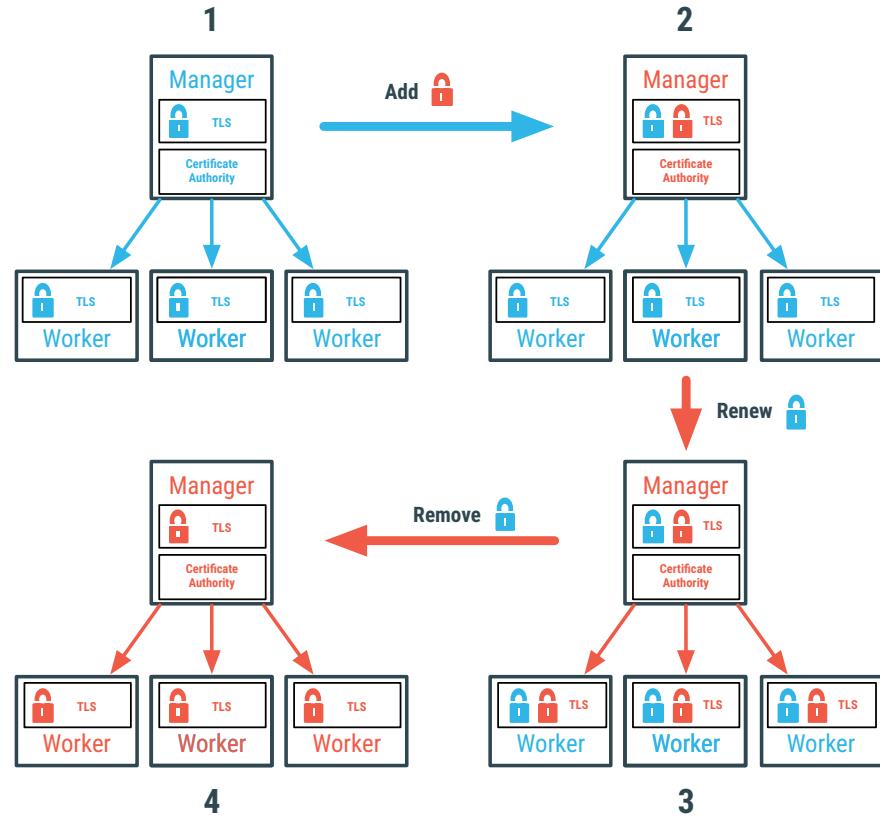
Transparent Root Rotation



3

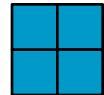
swarmKit

Transparent Root Rotation

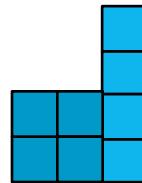


Bringing it all together

Notary for Docker image name resolution



Notary for Docker image name resolution



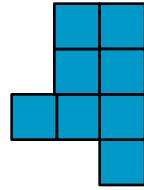
Cryptographically Verified Pulls



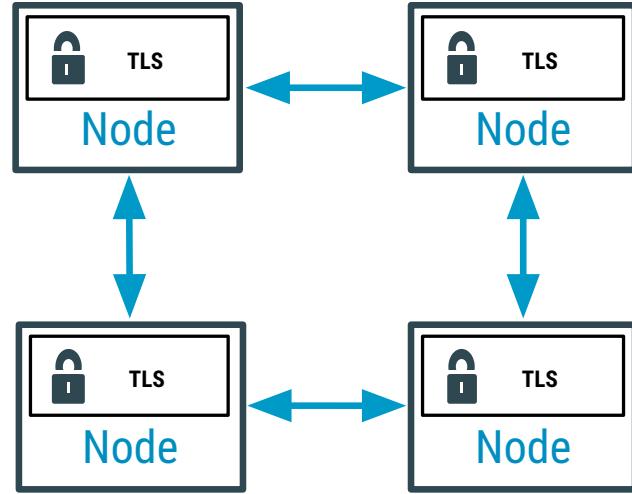
swarmKit delivered Docker containers



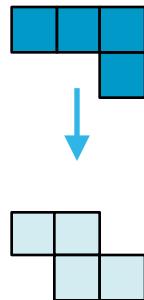
swarmKit delivered Docker containers



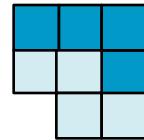
Authorized,
Authenticated,
Encrypted delivery
of Resources



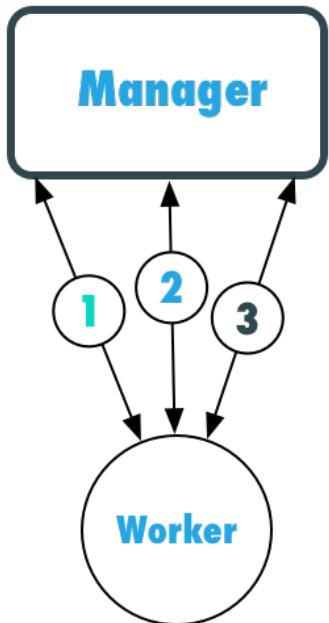
infraKit for swarmKit Bootstrap



infraKit for swarmKit Bootstrap

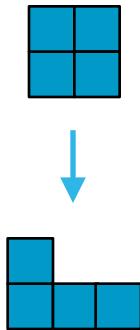


Secure Node Cluster Introduction

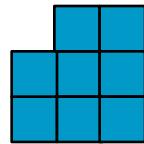


1. Retrieve and validate Root CA Public key material.
2. Submit new CSR along with secret token.
3. Retrieve the signed certificate.

linuxKit as the base OS builder



linuxKit as the base OS builder



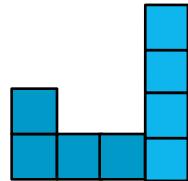
Hardened Configuration



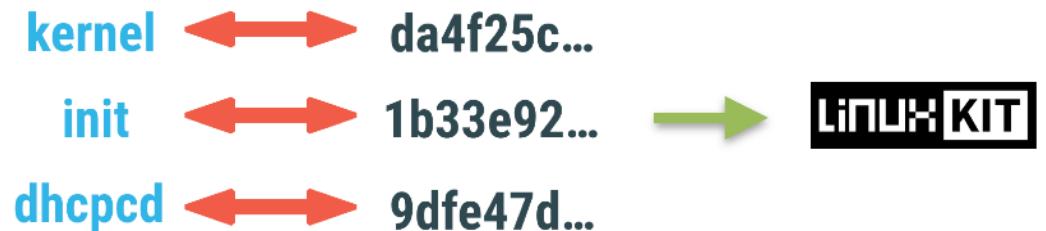
Notary for secure dependency resolution



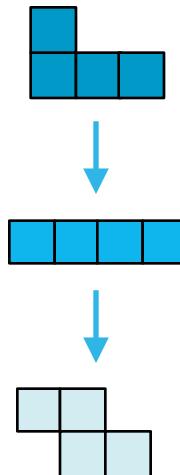
Notary for secure dependency resolution



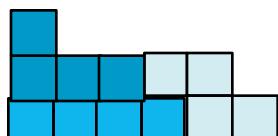
Cryptographically Verified Build



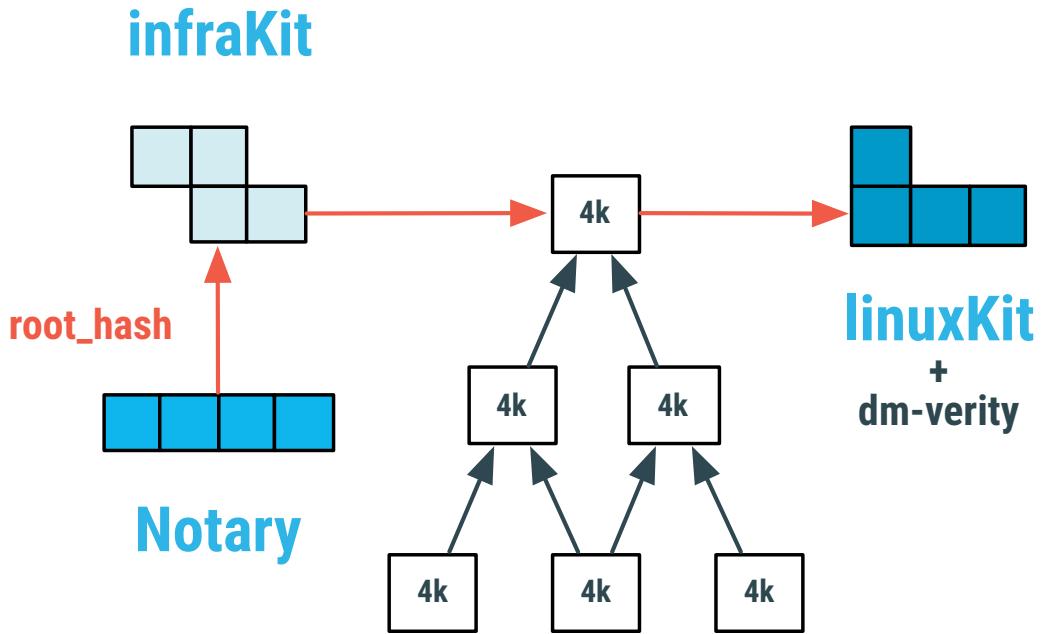
infraKit plus Notary for trusted OS Provisioning

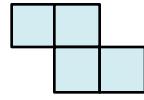


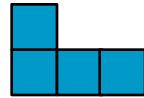
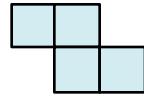
infraKit plus Notary for trusted OS Provisioning

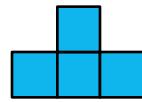
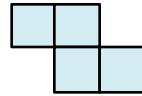


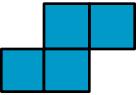
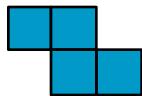
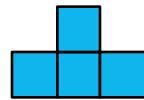
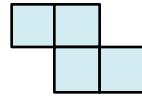
Cryptographically Verified Boot

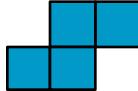
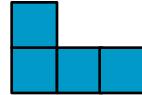
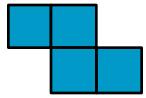
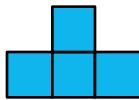
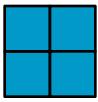
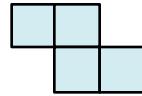


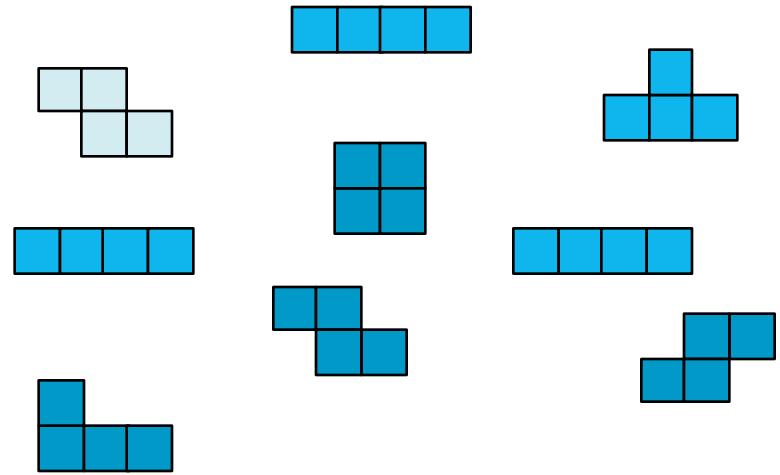


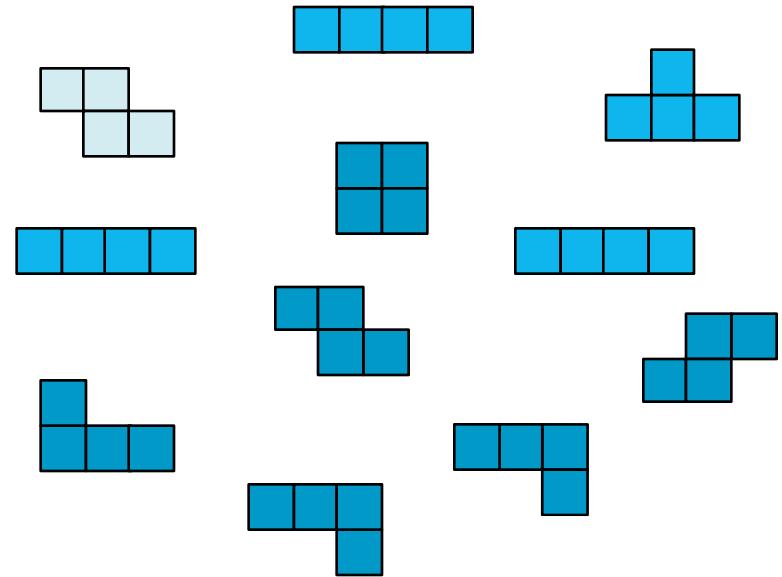


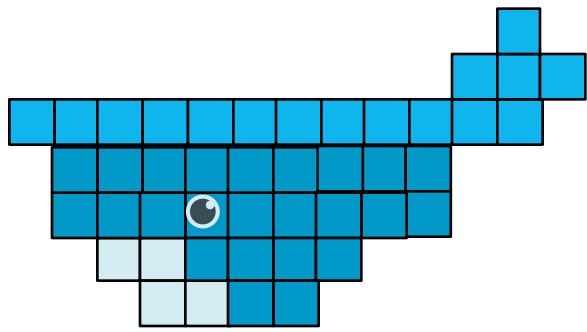


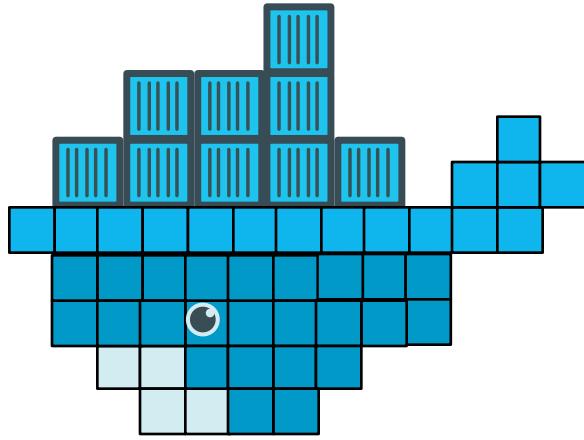












Thank you!